

Files and File Systems

Major OS Themes

Virtualization

- Present physical resource as a more general, powerful, or easy-to-use form of itself
- Present illusion of multiple (or unlimited) resources where only one (or a few) really exist
- Examples: CPU, Memory, **Hard drives**

Concurrency

- Coordinate multiple activities to ensure correctness

Persistence

- Some data needs to survive crashes and power failures
- Need abstractions, mechanisms, policies for all

How to virtualize persistent storage?

- Recall OS Goals:
 - **Convenience** for the user and **efficient** use of machine
- Virtualization and abstraction helps with first goal
 - Files and directories abstract away the hard drive
- Efficiency:
 - File system controls when and how data is transferred to persistent storage

Outline

- Persistent storage of data
- Files
- Directories
- File Operations
- How to specify which disk blocks belong to a file

File Systems

- Provide long-term information storage
- Requirements:
 - Store very large amounts of information
 - Information must survive the termination of process using it
 - Multiple processes must be able to access info concurrently
- Two views of file systems:
 - User view – convenient logical organization of information
 - OS view – managing physical storage media, enforcing access restrictions

File Systems

- Implement an abstraction ([files](#)) for secondary storage
- Organize files logically ([directories](#))
- Permit sharing of data between processes, people, and machines
- Protect data from unwanted access (security)

File Operations


- Creation
 - Find space in file system, add entry to **directory**
 - map **file name** to **location and attributes**
- Writing
- Reading
 - Dominant abstraction is “**file as stream**”
- Repositioning within a file
- Deleting a file
- Truncation and appending
 - May erase the contents (or part of the contents) of a file while keeping attributes

File Access Methods

- General-purpose file systems support simple methods
 - **Sequential access** – read bytes one at a time, in order
 - **Direct access** – random access given block/byte number
- Database systems support more sophisticated methods
 - **Record access** - fixed or variable length
 - **Indexed access**
- What file access method(s) does Unix/Linux, Windows provide?

Directories

- Directories provide **logical structure** to file systems
 - For users, they provide a means to **organize files**
 - For the file system, they provide a **convenient naming interface**
 - Separates logical file organization from physical file placement
 - Stores information about files (owner, permission, etc.)
- Most file systems support multi-level directories
 - Naming hierarchies (`/`, `/usr`, `/usr/local/`, `/home`, ...)




Directory Structure

- A directory is a list of entries – names and associated metadata
 - Metadata is not the data itself, but information that describes properties of the data (size, protection, location, etc.)
- List is usually unordered (effectively random)
 - Entries usually sorted by program that reads directory
- Directories typically stored in files
 - Only need to manage one kind of secondary storage unit


Directory Implementations

- Single-level, two-level, or tree-structured
- Acyclic-graph directories: allows for shared directories
 - The same file or subdirectory may be in 2 different directories

Tree-structured:



Acyclic graph:




Directory Implementation

- Option 1: List
 - Simple list of file names and pointers to file metadata
 - Requires linear search to find entries
 - Easy to implement, slow to execute
 - And **directory operations are frequent!**
- Option 2: Hash Table
 - Create a list of file info structures
 - Hash file name to get a pointer to the file info structure in the list
 - Hash table takes space


Hard Links

- In this example there is one hard link to **local**




Hard Links

- Except that there really are 3 links to local because each subdirectory has a link to its parent (“..”)




Hard Links

- The file **afile** has one hard link to it.



Hard Links


- We can create another hard link to the file using the ln command: `ln /usr/afile /bfile`



There is only one file, but it has two pointers (aliases) to it with potentially different names.

Hard Links


- If we remove **afile**, nothing happens to **bfile**.



Symbolic link

- Let's create a symbolic link instead:

```
ln -s /usr/afile /bfile
```




Instead of creating a new hard link, a special kind of file is created that contains the path to the linked file.

The file system code knows to use the path to find the contents.

Symbolic link

- Now what happens if we remove **afile**?



afile is removed and we have a “dangling pointer”

Hard links vs Symbolic Links

Hard Links

- Can't create new hard links to directories (might create cycles)
- Can't create hard links across partitions.
- No extra processing to follow hard links.
- Removing a hard link only removes the file if it is the last link to the file.


Symbolic Links

- Can create symbolic links to directories
- Can create symbolic links across partitions
- File system needs to look up link to follow it
- Removing a file may lead to a dangling link.

Next Up

OS views a disk as an array of fixed-size blocks.


- Implementing File Systems
 - Disk layout
 - File metadata
 - Directory data



Associating Data Blocks to Files

- We need a strategy to store which disk blocks belong to a file.
- Possibilities:
 - Contiguous Allocation
 - Linked Allocation (FAT)
 - Indexed Allocation (FFS, EXT2/3/4 — Inodes)
 - Extent-based Allocation (NTFS, Assignment 1)
- We also need somewhere to store the file metadata
 - permissions, last modified time, owner, type, ...


Contiguous Allocation



File Name	Start Blk	Length
File A	2	3
File B	9	5
File C	18	8
File D	27	2

Each file is given a contiguous set of blocks.
Store the starting block and length in the directory.

Linked Allocation

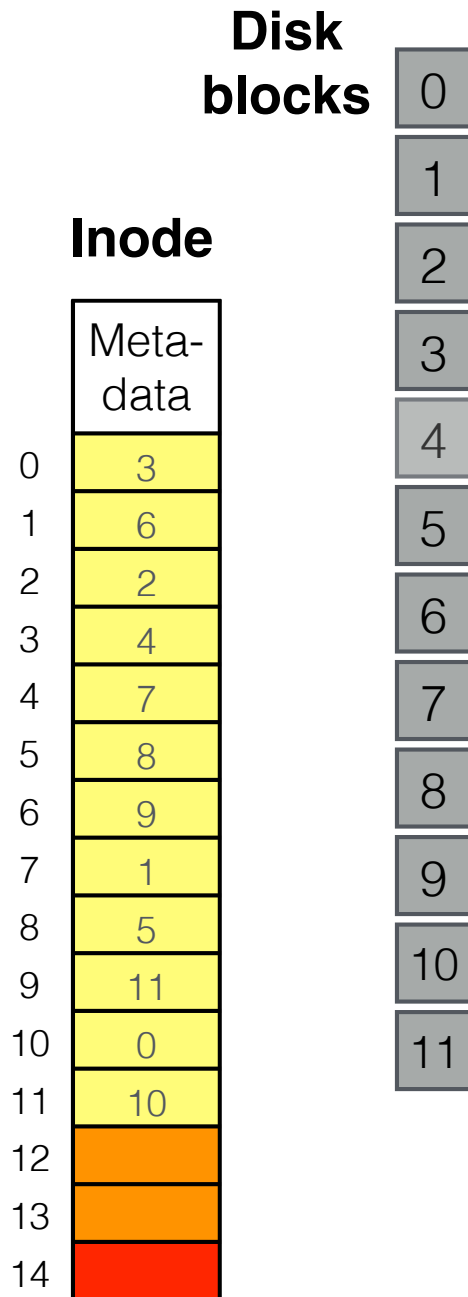


File Name	Start Blk	Last Blk
...
File B	1	22
...

Store the starting block and length in the directory.
Each block in a file contains a pointer to the next block
FAT is similar — We will discuss later

Indexed Allocation (Inodes)

- Unix **inodes** implement an indexed structure for files
- All file metadata is stored in an inode
 - Unix directory entries map file names to inodes
- Each inode contains 15 block pointers
 - First 12 are direct block pointers
 - Disk addresses of first 12 data blocks in file
 - The 13th is a single indirect block pointer
 - Address of block containing addresses of data blocks
 - Then the 14th is a double indirect block pointer
 - Address of block containing addresses of single indirect blocks
 - Finally, the 15th is a triple indirect block pointer




Inodes

- Each file and directory is represented by an inode
- First 12 pointers point directly to data blocks
 - 0th pointer points to the first block of the file
 - 11th pointer points to the 12th block of the file.
- The blocks do not need to be contiguous.

Remember, a disk is viewed as an array of blocks. So the “pointer” is really an index into the array of disk blocks

Inodes




- Pointer 12 points to a disk block that contains pointers to data blocks.

Inodes

	Inode
	Meta-data
0	3
1	6
2	2
3	4
4	7
5	8
6	9
7	1
8	5
9	11
10	0
11	10
12	
13	
14	

- Pointer 13 points to a disk block that contains pointers to single indirect blocks.




Inodes

Inode


	Meta-data
0	3
1	6
2	2
3	4
4	7
5	8
6	9
7	1
8	5
9	11
10	0
11	10
12	
13	
14	

- Pointer 14 points to a disk block that contains pointers to double indirect blocks.



Putting it all together

- Ext2 Linux file system Inodes are 128 bytes



Disk Layout Strategies

- Files often span multiple disk blocks
- How do you find all of the blocks for a file?
 1. **Contiguous allocation**
 - Fast, simplifies directory access and allows indexing
 - Inflexible, causes external fragmentation, requires compaction
 2. **Linked**, or chained, structure
 - Each block points to the next, directory points to the first
 - Good for sequential (streaming) access, bad for all others
 3. **Indexed** structure (kind of like address translation)
 - An “index block” contains pointers to many other blocks
 - Handles random access better, still good for sequential
 - May require multiple, linked index blocks