

CSC373

Algorithm Design, Analysis & Complexity

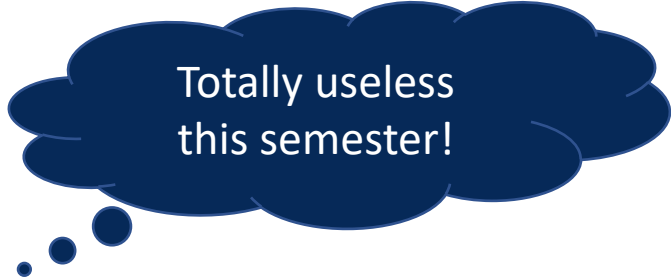
Nisarg Shah

Introduction

- **Instructors**

- **Nisarg Shah**

- cs.toronto.edu/~nisarg, nisarg@cs, SF 2301C
 - LEC 0101 and 0102



Totally useless
this semester!

- **TAs:** Too many to list

- **Disclaimer!**

- First online version of the course, so expect a bumpy ride at the start, but hopefully, we'll get through together
 - Use any of the feedback mediums (email, Piazza, ...) to let me know if you have any suggestions for improvement

Course Information

- **Course Page** www.cs.toronto.edu/~nisarg/teaching/373f20/
 - All the information below is in the course information sheet, available on Piazza
- **Discussion Board** piazza.com/utoronto.ca/fall2020/csc373
- **Grading – MarkUs**
 - Link will be distributed after about a week or two
 - LaTeX preferred, scans are OK!
- All times in **Eastern time zone**, all zoom links on the course page

Lectures

- **Time & Place:** Tue 4-5pm, Thu 1-3pm, Zoom
- **Details**
 - Delivered live
 - 10 minute break after every 50 minutes of lecture
 - Students can ask questions using Zoom's chat feature
 - One TA will be present to continuously answer questions
 - I might also answer questions once in a while

Tutorials

- **Time & Place:** Tue 5-6pm, Zoom
- **Details**
 - Delivered live by TAs
 - Problem sets will be posted early on the course webpage
 - Easier problems that are warm-up to assignments/exams
 - Please try them before coming to the tutorials
 - TAs will explain the problems, allow you to discuss them in breakout rooms, and then go over key parts of the solutions
 - Solutions will be posted later on the course webpage

Tutorials

- Further details

- Each section is divided into three parts (A,B,C)
- Students divided by birth month: A = Jan-Apr, B = May-Aug, C = Sep-Dec
- Feel free to attend a different tutorial than the one you're assigned
 - EXCEPT when the tutorial slot is being used for a test
- If the attendance is low, the number of tutorials per section may be reduced

Office Hours

- **Time & Place:** Wed 4-5pm, Fri 10-11am, Zoom
 - Do you have conflicts with these slots? **Poll!**
- **Details**
 - I will conduct them
 - Use the “raise hand” feature
 - I will call upon the raised hands in order
 - When called upon, unmute and ask the question
 - Always phrase your question in a way that doesn’t give away your solutions or approach to an assignment problem
 - Just like in a physical office

Tests

- 2 term tests, one end-of-term test (final exam)
- **Time & Place:** Tue 5-6pm (tutorial slot)
 - Need to be able to attend live!
 - I'm considering using part of the Tue 4-5pm lecture slot to give you more time
- **Tentative Plan**
 - Open book, closed internet
 - You may be asked to join a zoom link and keep your video on
 - If you have a question, you can “raise hand”, and I or a TA can take you to a breakout room to answer your question
 - Upload scanned answer sheet at the end (we'll do a mock run of this)

Assignments

- 4 assignments, best 3 out of 4
- Group work
 - In groups of up to three students
 - Best way to learn is for each member to try each problem
- Questions will be more difficult
 - May need to mull them over for several days; do *not* expect to start and finish the assignment on the same day!
 - May include bonus questions
- Submission on MarkUs, more details later
 - May need to compress the PDF

Grading Policy

- 3 homeworks * 10% = 30%
- 2 term tests * 20% = 40%
- Final exam * 30% = 30%

- **NOTE:** To pass, you must earn at least 40% on the final exam

Approximate Due Dates

- Please note the word **approximate**!
 - Assignment 1: Apx. Oct 9
 - Assignment 2: Apx. Oct 30
 - Assignment 3: Apx. Nov 13
 - Assignment 4: Apx. Nov 27
 - Midterm 1: Apx. Oct 20
 - Midterm 2: Apx. Nov 17
- Conflicts
 - The tests are during the tutorial slot, so there should ideally be no conflict
 - That said, if you think you'll have a conflict, let me know at the earliest

Textbook

- Primary reference: lecture slides
- Primary textbook (required)
 - [CLRS] Cormen, Leiserson, Rivest, Stein: *Introduction to Algorithms*.
- Supplementary textbooks (optional)
 - [DPV] Dasgupta, Papadimitriou, Vazirani: *Algorithms*.
 - [KT] Kleinberg; Tardos: *Algorithm Design*.

Other Policies

- Collaboration

- Free to discuss with classmates or read online material
- Must write solutions in your own words
 - Easier if you do not take any pictures/notes from discussions

- Citation

- For each question, must cite the peer (write the name) or the online sources (provide links), if you obtained a significant insight directly pertinent to the question
- Failing to do this is plagiarism!

Other Policies

- “No Garbage” Policy

- Borrowed from: Prof. Allan Borodin (citation!)

1. Partial marks for viable approaches
2. Zero marks if the answer makes no sense
3. 20% marks if you admit to not knowing how to approach the question (“I do not know how to approach this question”)

- $20\% > 0\%$!!

Other Policies

- Late Days

- 4 total late days across all 4 assignments
- Managed by MarkUs
- At most 2 late days can be applied to a single assignment
- Already covers legitimate reasons such as illness, university activities, etc.
 - Petitions will only be granted for circumstances which cannot be covered by this

Zoom Features

- Just to get acquainted, let's try out the following features:
 - Polls (already tried)
 - Chat
 - Reactions
 - Raise hand
 - Yes/No
 - Breakout rooms

Enough with the
boring stuff.

What will we study?

Why will we study it?



Muhammad ibn Musa al-Khwarizmi
c. 780 – c. 850

What is this course about?

- Algorithms

- Ubiquitous in the real world
 - From your smartphone to self-driving cars
 - From graph problems to graphics problems
 - ...
- Important to be able to design and analyze algorithms
- For some problems, good algorithms are hard to find
 - For some of these problems, we can formally establish complexity results
 - We'll often find that one problem is easy, but its minor variants are suddenly hard

What is this course about?

- Algorithms

- Algorithms in specialized environments or using advanced techniques
 - Distributed, parallel, streaming, sublinear time, spectral, genetic...
- Other concerns with algorithms
 - Fairness, ethics, ...
- ...mostly beyond the scope of this course

What is this course about?

- **Topics in this course**

- Divide and Conquer
- Greedy
- Dynamic programming
- Network flow
- Linear programming
- NP-completeness (not really an algorithm design paradigm)
- Approximation algorithms (if time permits)
- Randomized algorithms (if time permits)

What is this course about?

- How do we know which paradigm is right for a given problem?
 - A very interesting question!
 - Subject of much ongoing research...
 - Sometimes, you just know it when you see it...
- How do we analyze an algorithm?
 - Proof of correctness
 - Proof of running time
 - We'll try to prove the algorithm is *efficient* in the *worst case*
 - In practice, average case matters just as much (or even more)

What is this course about?

- What does it mean for an algorithm to be efficient in the worst case?
 - Polynomial time
 - It should use at most $\text{poly}(n)$ steps on *any* n -bit input
 - $n, n^2, n^{100}, 100n^6 + 237n^2 + 432, \dots$
 - If the input to an algorithm is a number x , the number of bits of input is $\log x$
 - This is because it takes $\log x$ bits to represent the input x in binary
 - So the running time should be polynomial in $\log x$, not in x
 - How much is too much?

What is this course about?

Picture-Hanging Puzzles*

Erik D. Demaine[†] Martin L. Demaine[†] Yair N. Minsky[‡] Joseph S. B. Mitchell[§]
Ronald L. Rivest[†] Mihai Pătraşcu[¶]

Theorem 7 *For any $n \geq k \geq 1$, there is a picture hanging on n nails, of length $n^{c'}$ for a constant c' , that falls upon the removal of any k of the nails.*

$n^{6,100 \log_2 c}$. Using the $c \leq 1,078$ upper bound, we obtain an upper bound of $c' \leq 6,575,800$. Using

So, while this construction is polynomial, it is a rather large polynomial. For small values of n , we can use known small sorting networks to obtain somewhat reasonable constructions.

What is this course about?

Better Balance by Being Biased: A 0.8776-Approximation for Max Bisection

Per Austrin^{*}, Siavosh Benabbas^{*}, and Konstantinos Georgiou[†]

has a lot of flexibility, indicating that further improvements may be possible. We remark that, while polynomial, the running time of the algorithm is somewhat abysmal; loose estimates places it somewhere around $O(n^{10^{100}})$; the running time of the algorithm of [RT12] is similar.

What is this course about?

- What if we can't find an efficient algorithm for a problem?
 - Try to prove that the problem is hard
 - Formally establish complexity results
 - NP-completeness, NP-hardness, ...
- We'll often find that one problem may be easy, but its simple variants may suddenly become hard
 - Minimum spanning tree (MST) vs bounded degree MST
 - 2-colorability vs 3-colorability

I'm not convinced.

Will I really ever need to
know how to design
abstract algorithms?

At the very least...

This will help you prepare for your
technical job interview!

Real Microsoft interview question:

- Given an array a , find indices (i, j) with the largest $j - i$ such that $a[j] > a[i]$
- Greedy? Divide & conquer?

Disclaimer

- The course is **theoretical in nature**
 - You'll be working with abstract notations, proving correctness of algorithms, analyzing the running time of algorithms, designing new algorithms, and proving complexity results.
- **Something for everyone...**
 - If you're somewhat scared going into the course
 - If you're already comfortable with the proofs, and want challenging problems

Related/Follow-up Courses

- **Direct follow-up**
 - CSC473: Advanced Algorithms
 - CSC438: Computability and Logic
 - CSC463: Computational Complexity and Computability
- **Algorithms in other contexts**
 - CSC304: Algorithmic Game Theory and Mechanism Design (self promotion!)
 - CSC384: Introduction to Artificial Intelligence
 - CSC436: Numerical Algorithms
 - CSC418: Computer Graphics

Divide & Conquer

History?

- Maybe you saw a subset of these algorithms?
 - Mergesort - $O(n \log n)$
 - Karatsuba algorithm for fast multiplication - $O(n^{\log_2 3})$ rather than $O(n^2)$
 - Largest subsequence sum in $O(n)$
 - ...
- Have you seen some divide & conquer algorithms before?
 - Maybe in CSC236/CSC240 and/or CSC263/CSC265
 - Write “yes”/”no” in chat

Divide & Conquer

- **General framework**
 - Break (a large chunk of) a problem into two smaller subproblems of the same type
 - Solve each subproblem recursively and independently
 - At the end, quickly combine solutions from the two subproblems and/or solve any remaining part of the original problem
- Hard to formally define when a given algorithm is divide-and-conquer...
- Let's see some examples!

Master Theorem

- Here's the master theorem, as it appears in CLRS
 - Useful for analyzing divide-and-conquer running time
 - If you haven't already seen it, please spend some time understanding it

Theorem 4.1 (Master theorem)

Let $a \geq 1$ and $b > 1$ be constants, let $f(n)$ be a function, and let $T(n)$ be defined on the nonnegative integers by the recurrence

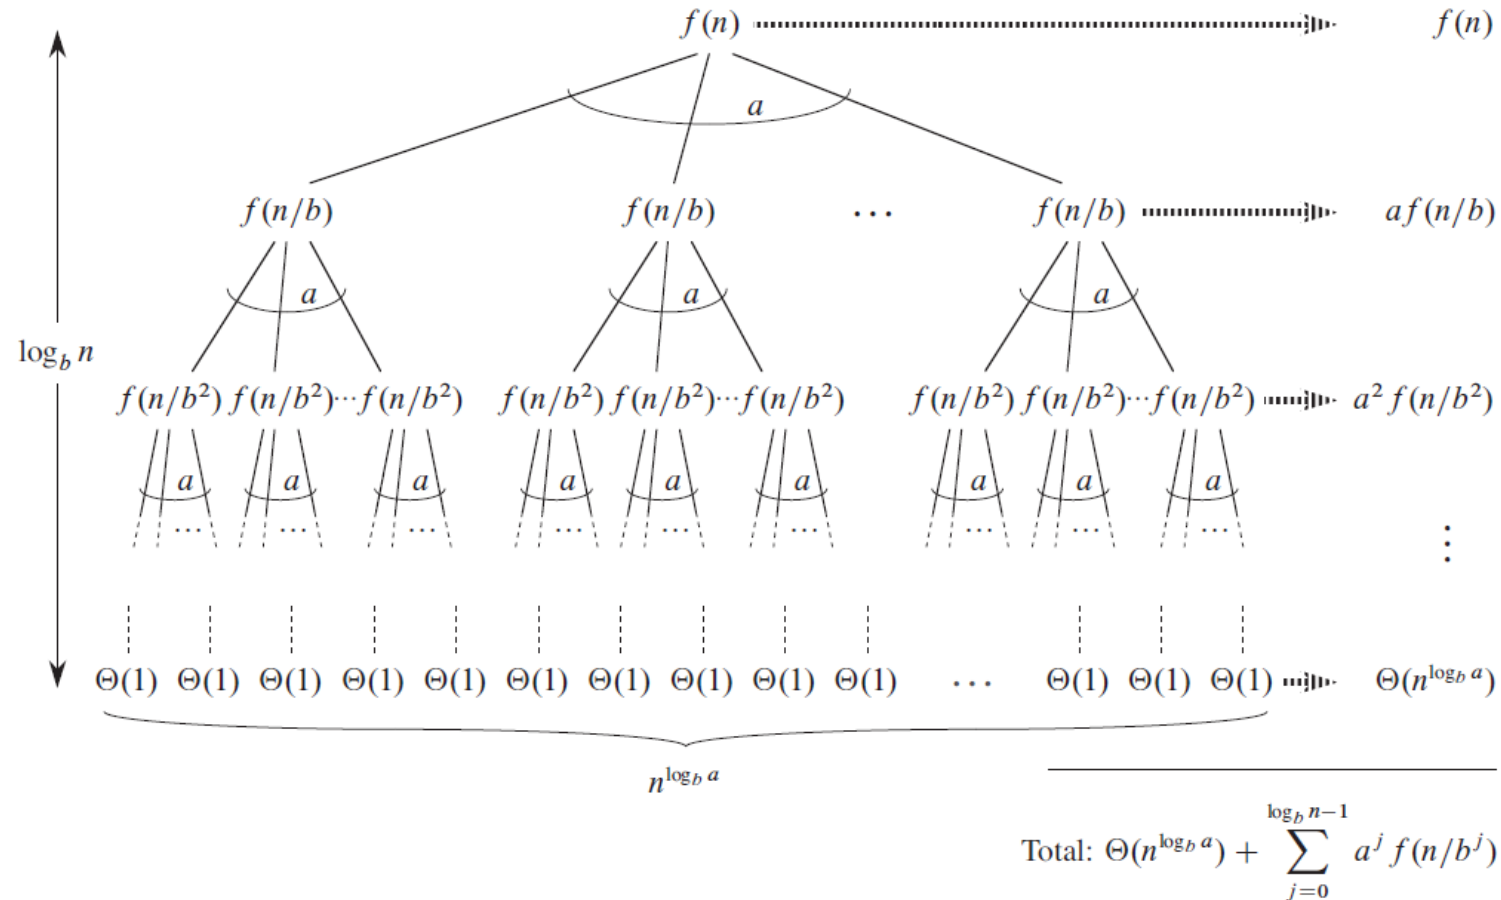
$$T(n) = aT(n/b) + f(n) ,$$

where we interpret n/b to mean either $\lfloor n/b \rfloor$ or $\lceil n/b \rceil$. Then $T(n)$ has the following asymptotic bounds:

1. If $f(n) = O(n^{\log_b a - \epsilon})$ for some constant $\epsilon > 0$, then $T(n) = \Theta(n^{\log_b a})$.
2. If $f(n) = \Theta(n^{\log_b a})$, then $T(n) = \Theta(n^{\log_b a} \lg n)$.
3. If $f(n) = \Omega(n^{\log_b a + \epsilon})$ for some constant $\epsilon > 0$, and if $af(n/b) \leq cf(n)$ for some constant $c < 1$ and all sufficiently large n , then $T(n) = \Theta(f(n))$. ■

Master Theorem

Intuition: Compare $f(n)$ with $n^{\log_b a}$. The larger determines the recurrence solution.



Counting Inversions

- Problem

- Given an array a of length n , count the number of pairs (i, j) such that $i < j$ but $a[i] > a[j]$

- Applications

- Voting theory
- Collaborative filtering
- Measuring the “sortedness” of an array
- Sensitivity analysis of Google's ranking function
- Rank aggregation for meta-searching on the Web
- Nonparametric statistics (e.g., Kendall's tau distance)

Counting Inversions

- Problem

- Count (i, j) such that $i < j$ but $a[i] > a[j]$

- Brute force

- Check all $\Theta(n^2)$ pairs

- Divide & conquer

- **Divide:** break array into two equal halves x and y
- **Conquer:** count inversions in each half recursively
- **Combine:**
 - Solve (we'll see how): count inversions with one entry in x and one in y
 - Merge: add all three counts

Counting Inversions

- From Kevin Wayne's slides

Sort-And-Count (L)

IF list L has one element

RETURN $(0, L)$.

DIVIDE the list into two halves A and B .

$(r_A, A) \leftarrow \text{Sort-And-Count}(A)$.

$(r_B, B) \leftarrow \text{Sort-And-Count}(B)$.

$(r_{AB}, L') \leftarrow \text{Merge-And-Count}(A, B)$.

RETURN $(r_A + r_B + r_{AB}, L')$.

Counting Inversions

input

1	5	4	8	10	2	6	9	3	7
---	---	---	---	----	---	---	---	---	---

count inversions in left half A

1	5	4	8	10
---	---	---	---	----

5-4

count inversions in right half B

2	6	9	3	7
---	---	---	---	---

6-3 9-3 9-7

count inversions (a, b) with $a \in A$ and $b \in B$

1	5	4	8	10
---	---	---	---	----

2	6	9	3	7
---	---	---	---	---

4-2 4-3 5-2 5-3 8-2 8-3 8-6 8-7 10-2 10-3 10-6 10-7 10-9

output $1 + 3 + 13 = 17$

Counting Inversions


Q. How to count inversions (a, b) with $a \in A$ and $b \in B$?

A. Easy if A and B are sorted!

Count inversions (a, b) with $a \in A$ and $b \in B$, assuming A and B are sorted.

- Scan A and B from left to right.
- Compare a_i and b_j .
- If $a_i < b_j$, then a_i is not inverted with any element left in B .
- If $a_i > b_j$, then b_j is inverted with every element left in A .
- Append smaller element to sorted list C .

count inversions (a, b) with $a \in A$ and $b \in B$



merge to form sorted list C




Counting Inversions

- How do we formally prove correctness?
 - Induction on n is usually very helpful
 - Allows you to assume correctness of subproblems
- Running time analysis
 - Suppose $T(n)$ is the running time for inputs of size n
 - Our algorithm satisfies $T(n) = 2 T(n/2) + O(n)$
 - Master theorem says this is $T(n) = O(n \log n)$

Without Master Theorem

Let's say $T(n) = 2 T(n/2) + 2n$



Overall: $2n \log n$

Closest Pair in \mathbb{R}^2

- **Problem:**

- Given n points of the form (x_i, y_i) in the plane, find the closest pair of points.

- **Applications:**

- Basic primitive in graphics and computer vision
- Geographic information systems, molecular modeling, air traffic control
- Special case of nearest neighbor


- **Brute force:** $\Theta(n^2)$

Intuition from 1D?

- In 1D, the problem would be easily $O(n \log n)$
 - Sort and check!
- **Sorting attempt in 2D**
 - Find closest points by x coordinate
 - Find closest points by y coordinate
- **Non-degeneracy assumption**
 - No two points have the same x or y coordinate


Intuition from 1D?

- **Sorting attempt in 2D**
 - Find closest points by x or y coordinate
 - Doesn't work!



Closest Pair in \mathbb{R}^2

- Let's try divide-and-conquer!
 - **Divide:** points in equal halves by drawing a vertical line L
 - **Conquer:** solve each half recursively
 - **Combine:** find closest pair with one point on each side of L
 - Return the best of 3 solutions




Seems like $\Omega(n^2)$ ☹

Closest Pair in \mathbb{R}^2


- Combine

- We can restrict our attention to points within δ of L on each side, where $\delta =$ best of the solutions in two halves



Closest Pair in \mathbb{R}^2

- Combine (let δ = best of solutions in two halves)
 - Only need to look at points within δ of L on each side,
 - Sort points on the strip by y coordinate
 - Only need to check each point with next 11 points in sorted list!




Why 11?

- Claim:

- If two points are at least 12 positions apart in the sorted list, their distance is at least δ

- Proof:

- No two points lie in the same $\delta/2 \times \delta/2$ box
- Two points that are more than two rows apart are at distance at least δ



Recap: Karatsuba's Algorithm

- Fast way to multiply two n digit integers x and y
- **Brute force:** $O(n^2)$ operations
- Karatsuba's observation:
 - Divide each integer into two parts
 - $x = x_1 * 10^{n/2} + x_2, y = y_1 * 10^{n/2} + y_2$
 - $xy = (x_1y_1) * 10^n + (x_1y_2 + x_2y_1) * 10^{n/2} + (x_2y_2)$
 - Four $n/2$ -digit multiplications can be replaced by three
 - $x_1y_2 + x_2y_1 = (x_1 + x_2)(y_1 + y_2) - x_1y_1 - x_2y_2$
 - Running time
 - $T(n) = 3 T(n/2) + O(n) \Rightarrow T(n) = O(n^{\log_2 3})$

[illegible]

Strassen's Algorithm

- Generalizes Karatsuba's insight to design a fast algorithm for multiplying two $n \times n$ matrices
 - Call n the “size” of the problem

$$\begin{bmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{bmatrix} = \begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} * \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix}$$

- Naively, this requires 8 multiplications of size $n/2$
 - $A_{11} * B_{11}, A_{12} * B_{21}, A_{11} * B_{12}, A_{12} * B_{22}, \dots$
- Strassen's insight: replace 8 multiplications by 7
 - Running time: $T(n) = 7 T(n/2) + O(n^2) \Rightarrow T(n) = O(n^{\log_2 7})$

Strassen's Algorithm

$$\begin{bmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{bmatrix} = \begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} * \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix}$$

```
STRASSEN(n, A, B)


---


IF (n = 1) RETURN A × B.
Partition A and B into 2-by-2 block matrices.
P1 ← STRASSEN(n / 2, A11, (B12 − B22)).
P2 ← STRASSEN(n / 2, (A11 + A12), B22).
P3 ← STRASSEN(n / 2, (A21 + A22), B11).
P4 ← STRASSEN(n / 2, A22, (B21 − B11)).
P5 ← STRASSEN(n / 2, (A11 + A22) × (B11 + B22)).
P6 ← STRASSEN(n / 2, (A12 − A22) × (B21 + B22)).
P7 ← STRASSEN(n / 2, (A11 − A21) × (B11 + B12)).
C11 = P5 + P4 − P2 + P6.
C12 = P1 + P2.
C21 = P3 + P4.
C22 = P1 + P5 − P3 − P7.
RETURN C.
```

assume *n* is a power of 2

keep track of indices of submatrices (don't copy matrix entries)

Median & Selection


- **Selection:**
 - Given array A of n comparable elements, find k th smallest
 - $k = 1$ is min, $k = n$ is max, $k = \lfloor (n + 1)/2 \rfloor$ is median
 - $O(n)$ is easy for min/max
- **What about k -selection?**
 - $O(nk)$ by modifying bubble sort
 - $O(n \log n)$ by sorting
 - $O(n + k \log n)$ using min-heap
 - $O(k + n \log k)$ using max-heap
- **Q:** What about just $O(n)$?
- **A:** Yes! Selection is easier than sorting.

QuickSelect

- Find a pivot p
- Divide A into two sub-arrays
 - A_{less} = elements $\leq p$, A_{more} = elements $> p$
 - If $|A_{less}| \geq k$, return k th smallest in A_{less} , otherwise return $(k - |A_{less}|)$ th smallest in A_{more}
- Problem?
 - If pivot is close to the min or the max, then we basically get $T(n) \leq T(n - 1) + O(n)$, which only gives $T(n) = O(n^2)$
 - Want to reduce $n - 1$ to a fraction of n (like $n/2$, $5n/6$, etc)

Finding a Good Pivot


- Divide n elements into $n/5$ groups of 5 each



N = 54

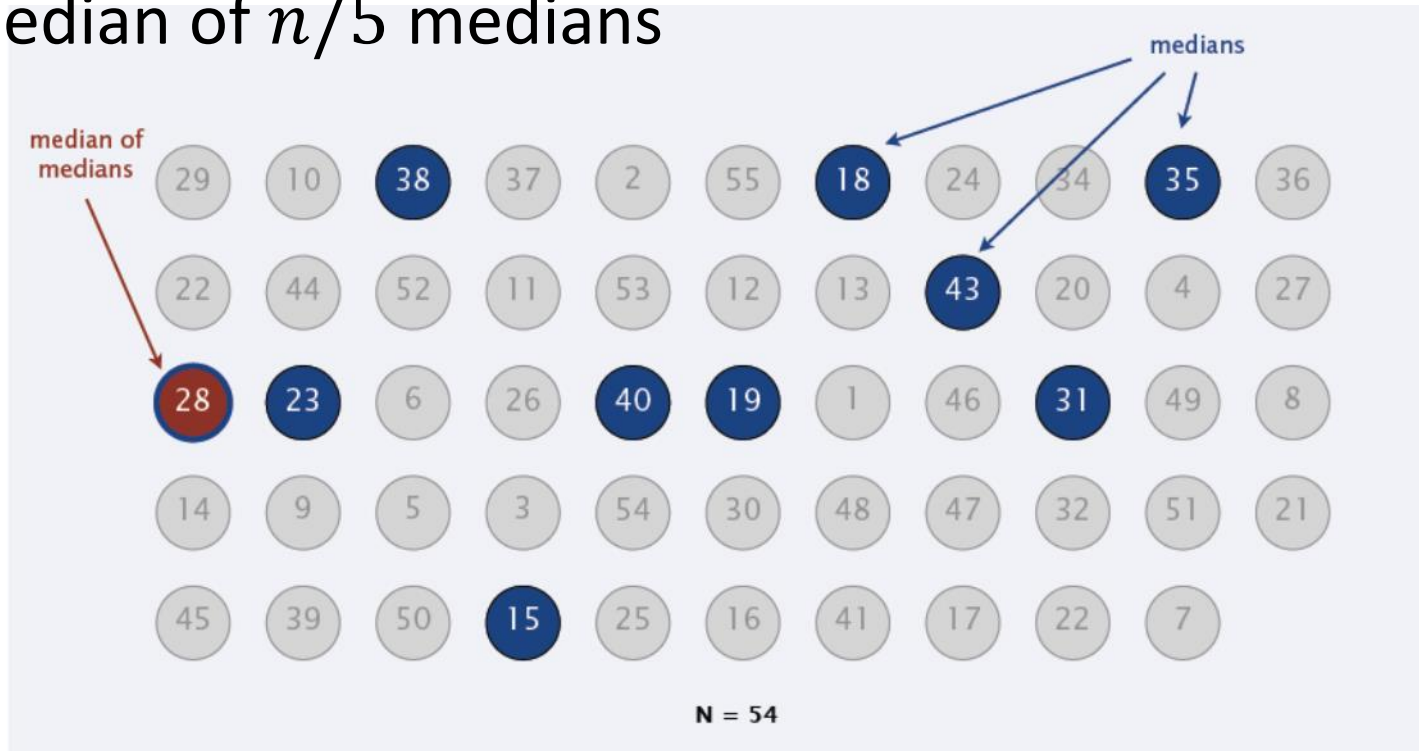
Finding a Good Pivot

- Divide n elements into $n/5$ groups of 5 each
- Find the median of each group



Finding a Good Pivot

- Divide n elements into $n/5$ groups of 5 each
- Find the median of each group
- Find the median of $n/5$ medians




Finding a Good Pivot

- Divide n elements into $n/5$ groups of 5 each
- Find the median of each group
- Find the median of $n/5$ medians
- Use this median of medians as the pivot in quickselect
- Q: Why does this work?

Analysis


- How many elements can be $\leq p^*$?
 - Out of $n/5$ medians, $n/10$ are $> p^*$



43

Analysis


- How many elements can be $\leq p^*$?
 - Out of $n/5$ medians, $n/10$ are $> p^*$



44

Analysis

- $n/10$ of the $n/5$ medians are $\leq p^*$
 - For each such median, there are 3 elements $\leq p^*$
 - So there can be at most $7n/10$ elements that can be $> p^*$




45

Analysis

- Thus, $|A_{more}| \leq 7^n/10$
 - Similarly, $|A_{less}| \leq 7^n/10$
 - (These are rough calculations...)
- How does this factor into overall algorithm analysis?

Analysis

- Divide n elements into $n/5$ groups of 5 each
- Find the median of each group
- Find p^* = median of $n/5$ medians
- Create A_{less} and A_{more} according to p^*
- Run selection on *one of* A_{less} or A_{more}



- $T(n) \leq T(n/5) + T(7n/10) + O(n)$
- Note: $n/5 + 7n/10 = 9n/10$
 - Only a fraction of n , so by the Master theorem, $T(n) = O(n)$

Residual Notes

- Best algorithm for a problem?
 - Typically hard to determine
 - We still don't know best algorithms for multiplying two n -digit integers or two $n \times n$ matrices
 - Integer multiplication
 - Breakthrough in March 2019: first $O(n \log n)$ time algorithm
 - It is conjectured that this is asymptotically optimal
 - Matrix multiplication
 - 1969 (Strassen): $O(n^{2.807})$
 - 1990: $O(n^{2.376})$
 - 2013: $O(n^{2.3729})$
 - 2014: $O(n^{2.3728639})$

Residual Notes

- Best algorithm for a problem?

- Usually, we design an algorithm and then analyze its running time

- Sometimes we can do the reverse:

- E.g., if you know you want an $O(n^2 \log n)$ algorithm

- Master theorem suggests that you can get it by

$$T(n) = 4 T(n/2) + O(n^2)$$

- So maybe you want to break your problem into 4 problems of size $n/2$ each, and then do $O(n^2)$ computation to combine

Residual Notes

- Access to input

- For much of this analysis, we are assuming random access to elements of input
- So we're ignoring underlying data structures (e.g. doubly linked list, binary tree, etc.)

- Machine operations

- We're only counting the number of comparison or arithmetic operations
- So we're ignoring issues like how real numbers are stored in the closest pair problem
- When we get to P vs NP, representation will matter

Residual Notes

- Size of the problem
 - Can be any reasonable parameter of the problem
 - E.g., for matrix multiplication, we used n as the size
 - But an input consists of two matrices with n^2 entries
 - It doesn't matter whether we call n or n^2 the size of the problem
 - The actual running time of the algorithm won't change