

## Search

- One of the most fundamental techniques in AI
  - Underlying sub-module in many AI systems
- Can solve many problems that humans are not good at.
- Can achieve super-human performance on some problems (Chess, go)
- Very useful as a general algorithmic technique for solving problems (both in AI and in other areas)

# How do we plan our holiday?

---

- We must take into account various preferences and constraints to develop a schedule.
- An important technique in developing such a schedule is “**hypothetical**” reasoning.
- Example: On holiday in England
  - Currently in Edinburgh
  - Flight leaves tomorrow from London
    - Need plan to get to your plane
    - If I take a 6 am train where will I be at 2 pm? Will I be still able to get to the airport on time?

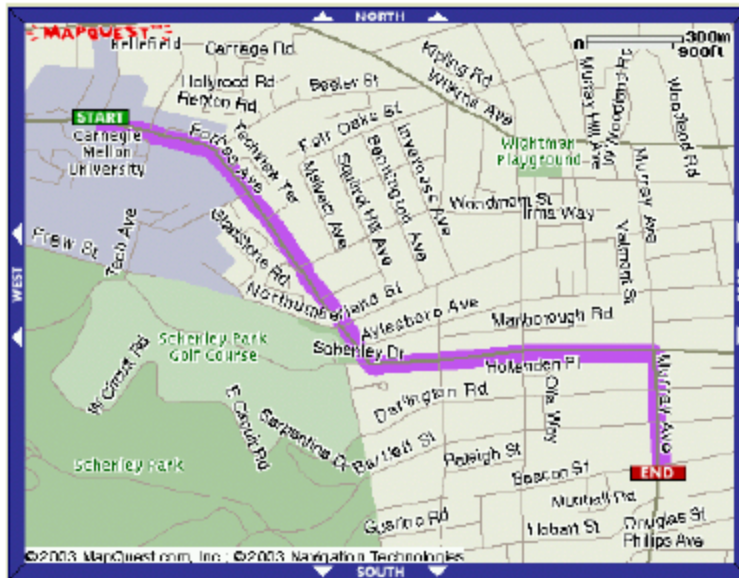
# How do we plan our holiday?

---

- This kind of hypothetical reasoning involves asking
  - what state will I be in after taking certain actions, or after certain sequences of events?
- From this we can reason about particular sequences of actions one should execute to achieve a desirable state.
- Search is a computational method for capturing a particular version of this kind of reasoning.

# Many problems can be solved by search:

## Search Problems



Slide 7

# Many problems can be solved by search:

---



Deepblue 1997

beats Kasparov world  
champion Chess player

AlphaGo 2016

beats Lee Sedol 9<sup>th</sup>  
dan Go player

2017 beats Ke Jie  
World #1 ranked player



# Why Search?

---

- Successful
  - Success in game playing programs based on search.
  - Many other AI problems can be successfully solved by search.
- Practical
  - Many problems don't have specific algorithms for solving them. Casting as search problems is often the easiest way of solving them.
  - Search can also be useful in approximation (e.g., local search in optimization problems).
  - Problem specific heuristics provides search with a way of exploiting extra knowledge.
- Some critical aspects of intelligent behavior, e.g., planning, can be naturally cast as search.

# Limitations of Search

---

- There are many difficult questions that are not resolved by search. In particular, the whole question of how does an intelligent system formulate the problem it wants to solve as a search problem is not addressed by search.
- Search only provides a method for solving the problem **once we have it correctly formulated.**

# Search

- Formulating a problem as search problem (representation)
- Heuristic Search
- Readings
  - Introduction: Chapter 3.1 – 3.3
  - Uninformed Search: Chapter 3.4
  - Heuristic Search: Chapters 3.5, 3.6



# Representing a problem: The Formalism

---

To model a problem as a search problem we need the following components:

1. **STATE SPACE:** A state is a representation of a configuration of the problem domain. The **state space** is a set of states included in our model of the problem.
2. **ACTIONS or STATE SPACE Transitions:** **Actions** these model the actions of the problem domain. In our model the domain actions are modeled as allowed transitions between state.

# Representing a problem: The Formalism

---

3. **INITIAL or START STATE and GOAL:** Identify the **initial state** that represents the starting conditions, and the goal state or goal condition one wants to achieve.
4. **Heuristics:** Formulate various **heuristics** to help guide the search process.

Note that this representation of the problem abstracts from the real problem (i.e., omits some details)—it is a model of the problem suitable for use in a computer.

Typically we only model relevant parts of the real world states, and the actions/initial and goal state/heuristics operate on that model in a way that reflects the real problem.

# The Formalism

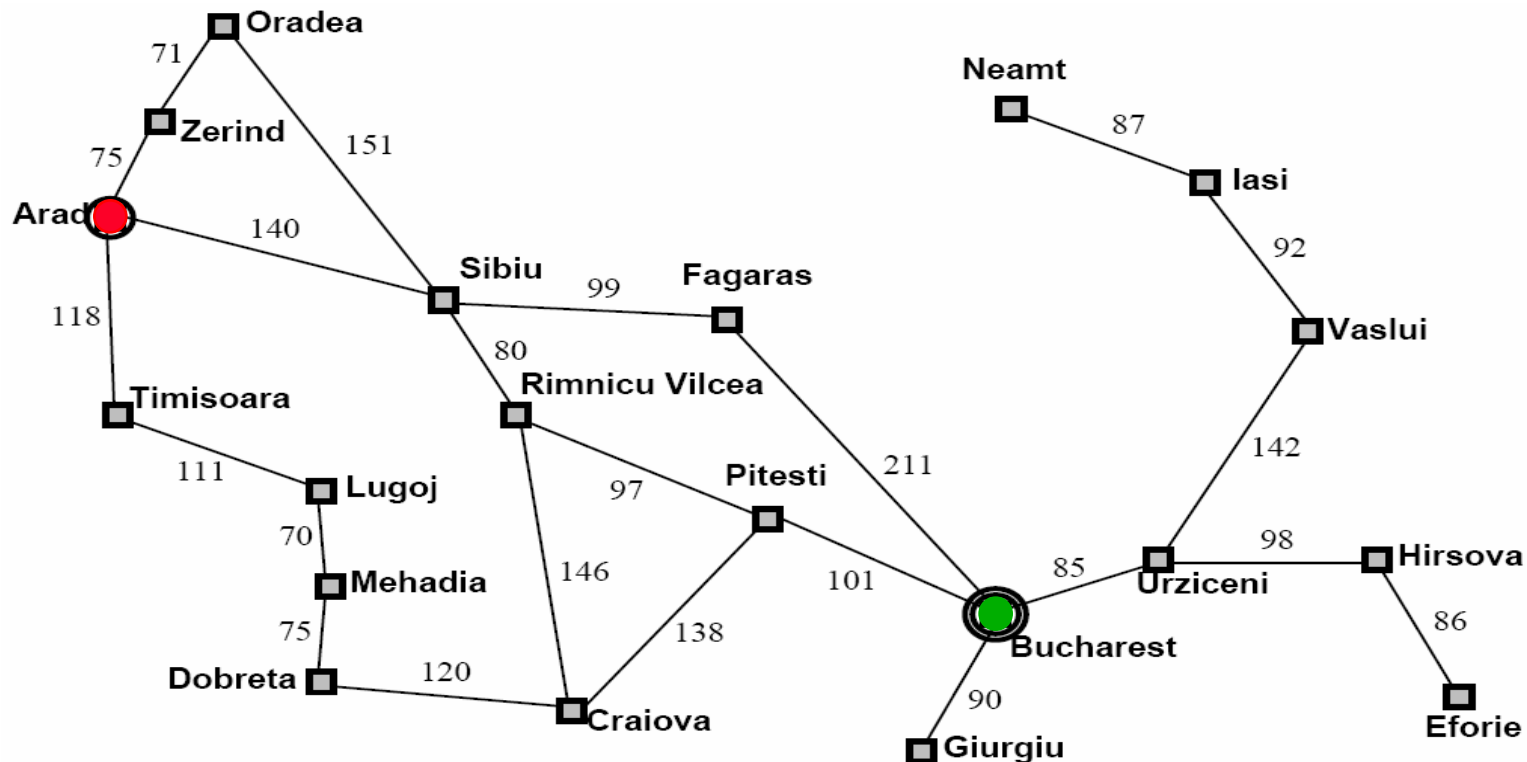
---

Once the problem has been formulated as a state space search, various algorithms can be utilized to solve the problem.

- A solution to the problem will be a sequence of actions/moves that can transform your current state into a state where your desired condition holds.

# Example 1: Romania Travel.

Currently in **Arad**, need to get to **Bucharest** by tomorrow to catch a flight. What is a reasonable **State Space**?



# Example 1. Romania Travel.

---

- State space.
  - **States**: The set of cities we can be in  $\{A, B, C, \dots, Z\}$ . E.g., the state A represents being in Arad.
    - Our abstraction: we are ignoring the low level details of driving, states where you are on the road between cities, etc.
  - **Actions**: drive between neighboring cities. This changes the state we are in (the city we are in).
  - **Initial state**: in Arad.
  - **Desired condition (Goal)**: be in a state where you are in Bucharest. (How many states satisfy this condition?)
- Solution will be the route, the sequence of cities to travel through to get to Bucharest.

# Example 2 Water Jugs.

---

- Water Jugs
  - We have a 3 gallon (liter) jug and a 4 gallon jug. We can fill either jug to the top from a tap, we can empty either jug, or we can pour one jug into the other (at least until the other jug is full).
  - **States**: each state is a pair of numbers (gal3, gal4)  
gal3 = the number of gallons in the 3 gallon jug  
gal4 = the number of gallons in the 4 gallon jug.
  - **Actions**: Empty-3-Gallon, Empty-4-Gallon, Fill-3-Gallon, Fill-4-Gallon, Pour-3-into-4, Pour 4-into-3.
  - **Initial state**: Various initial states are possible, e.g., (0,0)
  - **Desired condition (Goal)**: Various, e.g., (0,2) or (\*, 3) where \* means we don't care.

# Example 2 Water Jugs.

---

- Water Jugs
  - If we start off with gal3 and gal4 as integer, can only reach integer values.
  - Some values, e.g., (1,2) are not reachable from some initial state, e.g., (0,0).
  - Some actions are no-ops. They do not change the state, e.g.,
    - $(0,0) \rightarrow \text{Empty-3-Gallon} \rightarrow (0,0)$

## Example 3. The 8-Puzzle

---

7	2	4
5		6
8	3	1

Start State

1	2	3
4	5	6
7	8	

Goal State

**Rule:** Can slide a tile into the blank spot.

Alternative view: move the blank spot around.



# Example 3. The 8-Puzzle

---

- State space.
  - **States**: The different configurations of the tiles. How many different states? We can represent these states in a matrix or a 9 element vector.
  - **Actions**: **Moving the blank** up, down, left, right. Can every action be performed in every state?
  - **Initial state**: Various initial states are possible, e.g., the state shown on previous slide.
  - **Desired condition (Goal)**: be in a state where the tiles are all in the positions shown on the previous slide.
- Solution will be a sequence of moves of the blank that transform the initial state to a goal state.

**Question:** we could have as actions the movements of the individual tiles. Would this be a better representation?

## Example 3. The 8-Puzzle

---

- Although there are  $9!$  different configurations of the tiles (362,880) in fact the state space is divided into two disjoint parts.
- Only when the blank is in the middle are all four actions possible.
- Our goal condition is satisfied by only a single state. But one could easily have a goal condition like
  - The 8 is in the upper left hand corner.
  - How many different states satisfy this goal?


## Example 4: Vacuum World

---

- In the previous examples, a state in the search space represented a particular state of the world.
- However, states need not map directly to world configurations. Instead, a state could map to **knowledge states**.
- A knowledge state is a **set** of world states (set of ground states)—every world state that you believe to be possible.
- If you know the exact state of the world your knowledge state is set containing only that world state.
- The facts you know are those facts that are true in every world state contained in your knowledge state.

## Example 4. Vacuum World

- We have a vacuum cleaner and two rooms.
- Each room may or may not be dirty.
- The vacuum cleaner can move **left** or **right** (*the action has no effect if there is no room to the right/left*).
- The vacuum cleaner can **suck**; this cleans the room (*even if the room was already clean*).




Physical states

# Example 4. Vacuum World

## Knowledge-level State Space

- Each knowledge state consists of a set of possible world states. The agent knows that it is in one of these world states, but doesn't know which.




Goal is to have all rooms clean.

# Example 4. Vacuum World

## Knowledge-level State Space

- Complete knowledge of the world: agent knows exactly which physical state it is in. Then the states in the agent's state space consist of single physical states.
- Start in {5}:  
    <right, suck>




Goal is to have all rooms clean.

# Example 4. Vacuum World


## Knowledge-level State Space

- knowledge states:  
Agent's knowledge states consist of *sets of world states*.
- E.g. starting in {1,2,3,4,5,6,7,8}, the agent doesn't have any knowledge of where it is.
- Nevertheless, the action sequence *<right, suck, left, suck>* achieves the goal.



Goal is to have all rooms clean.

## Example 4. Vacuum World



Initial state.


{1,2,3,4,5,6,7,8}

Left

What does the agent know in this knowledge state?




## Example 4. Vacuum World




## Example 4. Vacuum World

---



## Example 4. Vacuum World

---





# Example 4. Vacuum World

---


- State space.
  - **States:**
    - ground states  $G = \{1, 2, 3, 4, 5, 6, 7, 8\}$ —the configuration shown on previous slide.
    - **states** =  $\{s \mid s \text{ is a non-empty subset of } G\}$ . There are  $2^8 - 1$ . Each state  $s$  is the set of world configurations the agent believes are possible. The agent does not know which ground state in  $s$  is the true world configuration.

# Example 4. Vacuum World

- **Actions:** Left, Right, Suck
  - How would you specify these actions?
  - $\text{Left}(s) = (s \cap \{1, 3, 5, 7\}) \cup \{g - 1 \mid g \in s \cap \{2, 4, 6, 8\}\}$
  - $\text{Right}(s) = (s \cap \{2, 4, 6, 8\}) \cup \{g + 1 \mid g \in s \cap \{1, 3, 5, 7\}\}$
  - $\text{Suck}(s) = \text{exercise}$
- e.g.,  $\text{Right}(\{5, 7\}) = \{6, 8\}$



Right



## Example 4. Vacuum World

---

- **Initial state**: Various initial states are possible, e.g., the state of not knowing anything about your room is  $\{1,2,3,4,5,6,7,8\}$ .
- **Desired condition (Goal)**:  $\{s \mid s \subseteq \{7, 8\}\}$  note this is a set of states (the goal need not be a single state)
- Solution will be a sequence of Left, Right, Suck moves that transform the initial state to a goal state.
- This example shows that the **state space** formalism can deal with domains in which our knowledge of the world state is incomplete.

# More complex situations

---

- Actions can lead to multiple states, e.g., flip a coin to obtain heads OR tails. Or we don't know for sure what the initial state is (prize is behind door 1, 2, or 3). Now we might want to consider how **likely** different states and action outcomes are.
- This leads to probabilistic models of the search space and different algorithms for solving the problem.
- Later we will see some techniques for reasoning under uncertainty.

# More complex situations

---

- The agent might be equipped with sensing actions.
  - These actions change the agent's knowledge state—they don't change the state of the world.
- With sensing we can search for contingent solutions: solutions that contain branches depending on the outcome of sensing actions.
  - <right, if dirt then suck, left>
- Searching for contingent plans needs different algorithms.



# Algorithms for Search

---

- AI search algorithms work with **implicitly defined** state spaces.
- There are typically an exponential number of states: impossible to explicitly represent them all.
- The space of possible configurations of a Go board is about  $3^{361}$  (standard 19X19 board).
- There are even more actions than state.

# Algorithms for Search

---

- In AI search we find solutions by constructing only those states we need to. In the worst case we will need to construct an exponential number of states—and the search might require too much computation.
- But often we can solve hard problems (like Go) while only examining a small fraction of the states.
- Hence the actions are given as compact successor state functions that when given a state  $x$  return the set of states  $S$  can be transformed to by the available actions.
  - This means that the state must contain enough information to allow the successor state function to perform its computation.

# Algorithms for Search

---

Inputs:

- a specified **initial state**  $I$
- a **successor** function  $S(x)$  yields the set of **all** states action pairs  $(y,a)$  such that state  $y$  can be reached from  $x$  by applying an action  $a$ . The successor function returns all states reachable by a single action from state  $x$ .
- a **goal test** a function  $G(x)$  that can be applied to a state  $x$  and returns true if  $x$  satisfies the goal condition.
- An **action cost** function  $C(x,a,y)$  which determines the cost of moving from state  $x$  to state  $y$  using action  $a$ . ( $C(x,a,y) = \infty$  if  $a$  does not yield  $y$  from  $x$ ). Note that different actions with different costs might generate the same transition  $x \rightarrow y$

# Algorithms for Search

---

## Output:

- a sequence of actions that transforms the initial state to a state satisfying the goal test.
  - Or just the sequence of states that arise from these actions (depends on what kind of information is most useful)
- The sequence might be, optimal in cost for some algorithms, optimal in length for some algorithms, come with no optimality guarantees from other algorithms.
  - That is, no other sequence transforms the initial state to a goal satisfying state with lower cost (or lesser length).

# Algorithms for Search

---

## Obtaining the action sequence.

- The set of successors of a state  $x$  might arise from different actions, e.g.,
  - $x \rightarrow a \rightarrow y$
  - $x \rightarrow b \rightarrow z$
- Successor function  $S(x)$  yields a set of states that can be reached from  $x$  via **any** action.
  - $S(x) = \{ \langle y, a \rangle, \langle z, b \rangle \}$   
 $y$  via action  $a$ ,  $z$  via action  $b$ .
  - $S(x) = \{ \langle y, a \rangle, \langle y, b \rangle \}$   
 $y$  via action  $a$ , also  $y$  via alternative action  $b$ .

# Search Algorithms

---

- The search space consists of **states** and actions that move between states.
- A **path** in the search space is a **sequence** of states connected by actions,  $\langle s_0, s_1, s_2, \dots, s_k \rangle$ , for every  $s_i$  and its successor  $s_{i+1}$  there must exist an action  $a_i$  that transitions  $s_i$  to  $s_{i+1}$ .
  - Alternately a path can be specified by
    - (a) an initial state  $s_0$ , and
    - (b) a sequence of actions that are applied in turn starting from  $s_0$ .
  - If the actions causing the transitions are specified then we can compute the cost of the path. E.g., if  $a_i$  is the action used to transition state  $s_i$  to  $s_{i+1}$  then the cost of the path is

$$\sum_{i=0}^{i=k-1} C(s_i, a_i, s_{i+1})$$

# Search Algorithms

---

- The search algorithms perform search by examining alternate **paths** of the search space.
  - The paths in the search space are extensions of previous paths, so to reduce space requirements paths can be stored as a **pair** containing the final state of the path and a pointer to the previous state. Following those pointers to the initial state yields the path.
  - For a path **p**: **p.final ()** is a function that returns the final state of the path **p**
- e.g.,  $p = \langle A, B, C \rangle$ ,  $p.\text{final}() = C$
- All paths examined by our search algorithms will start at the initial state I.

# Algorithm for Search

---

- We maintain a set of paths called the **OPEN** set (or frontier).
  - All of these paths start at the initial state. And they are **open** because we have not yet determined whether or not these paths are prefixes of a path leading to the a goal satisfying state.
- Initially we set  $OPEN = \{ \langle \text{Start State} \rangle \}$ 
  - The path hat starts and terminates at the start state.
- At each step we select a path **p** from OPEN.
  - We check if **p.final()** satisfies the goal,
  - If not we add all extensions of **p** to OPEN
    - for all successor states  $y \in S(p.final())$   
create a new path  $p_y = \langle p, y \rangle$  – extend the path p to include a transition to the state y
- e.g., if  $p = \langle a, b, c, d \rangle$  then  $p.final() = d$ . Let  $S(d) = \{e, f\}$  (i.e., states e and f can be reached from d via a single action)  
then we get two new extensions of p that will be added to OPEN
  1.  $p_e = \langle a, b, c, d, e \rangle$
  2.  $p_f = \langle a, b, c, d, f \rangle$



# Algorithm for Search

---

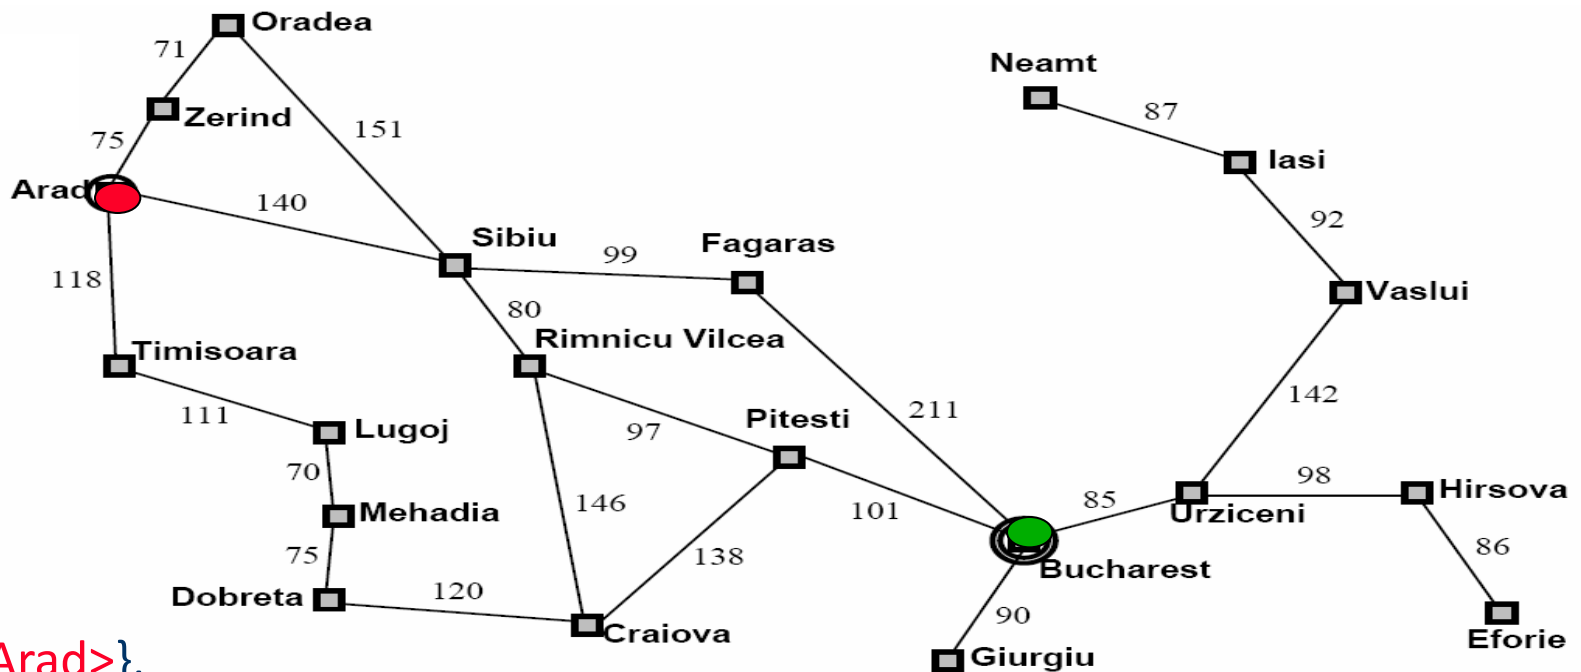
```
Search(open, S, goal?):
  open.insert(<start>)
  while not open.empty():
    p = open.extract()  #remove path from OPEN
    if (goal?(p.final())):
      return p          #p is a path reaching a goal satisfying
                        #state—it is a solution
    for succ in S(p.final()): #S is the successor function
      open.insert(<p,succ>)
                        #open could grow or shrink
  return false
```

When does OPEN get smaller in size?

# Algorithm for Search

---

- When a path  $p$  is extracted from open, we say that the algorithm **expands**  $p$ .
- The number of states we actually construct (i.e., the total number of states returned by the successor function  $S()$  summed over all calls to  $S()$ ), we hope is low compared to the total number of states in the search space
- The number of states expanded depends on the order of paths we extract from open.



{<Arad>},

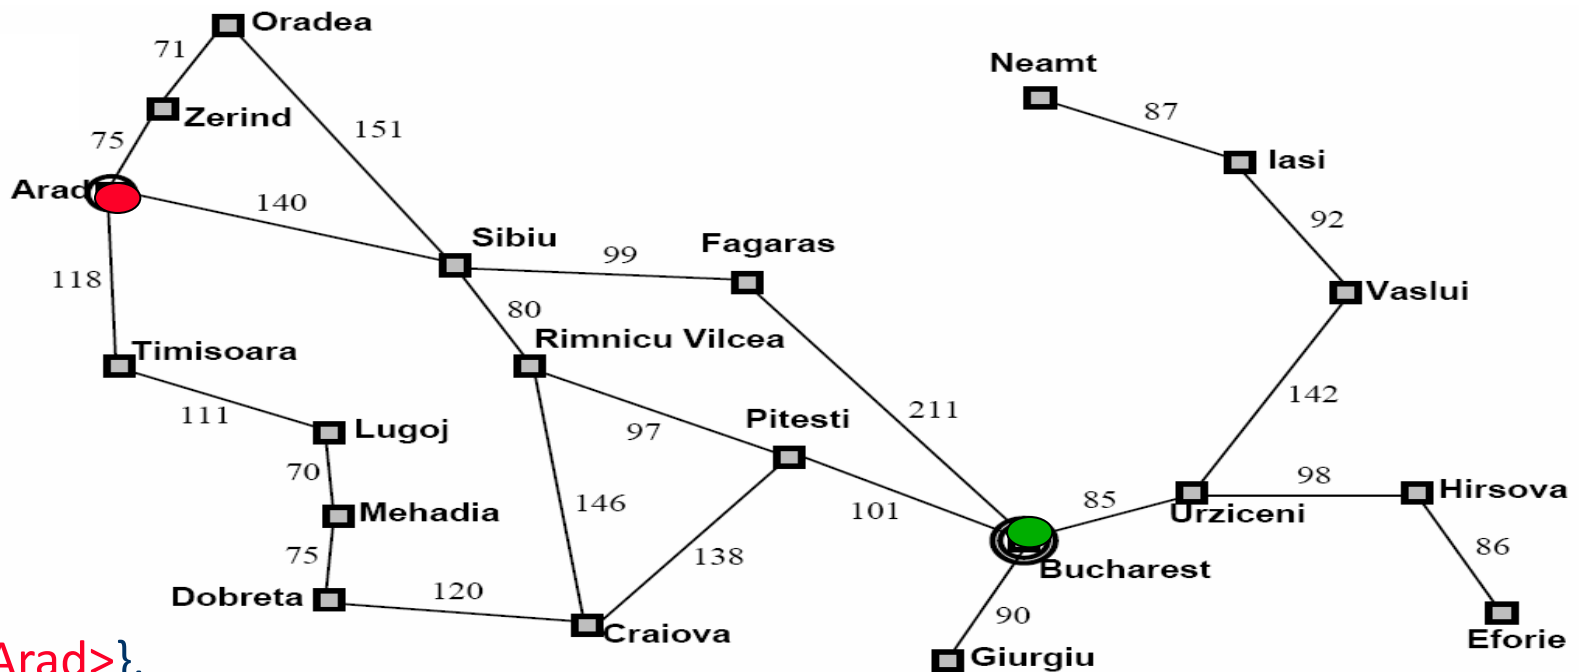
{<A,Z>, <A,T>, <A, S>},

{<A,Z>, <A,T>, <A,S,A>, <A,S,O>, <A,S,F>, <A,S,R>}

{<A,Z>, <A,T>, <A,S,A>, <A,S,O>, <A,S,R>, <A,S,F,S>, <A,S,F,B>}

**Solution:** Arad -> Sibiu -> Fagaras -> Bucharest

**Cost:** 140 + 99 + 211 = 450



{<Arad>},

{<A,Z>, <A,T>, <A,S>},

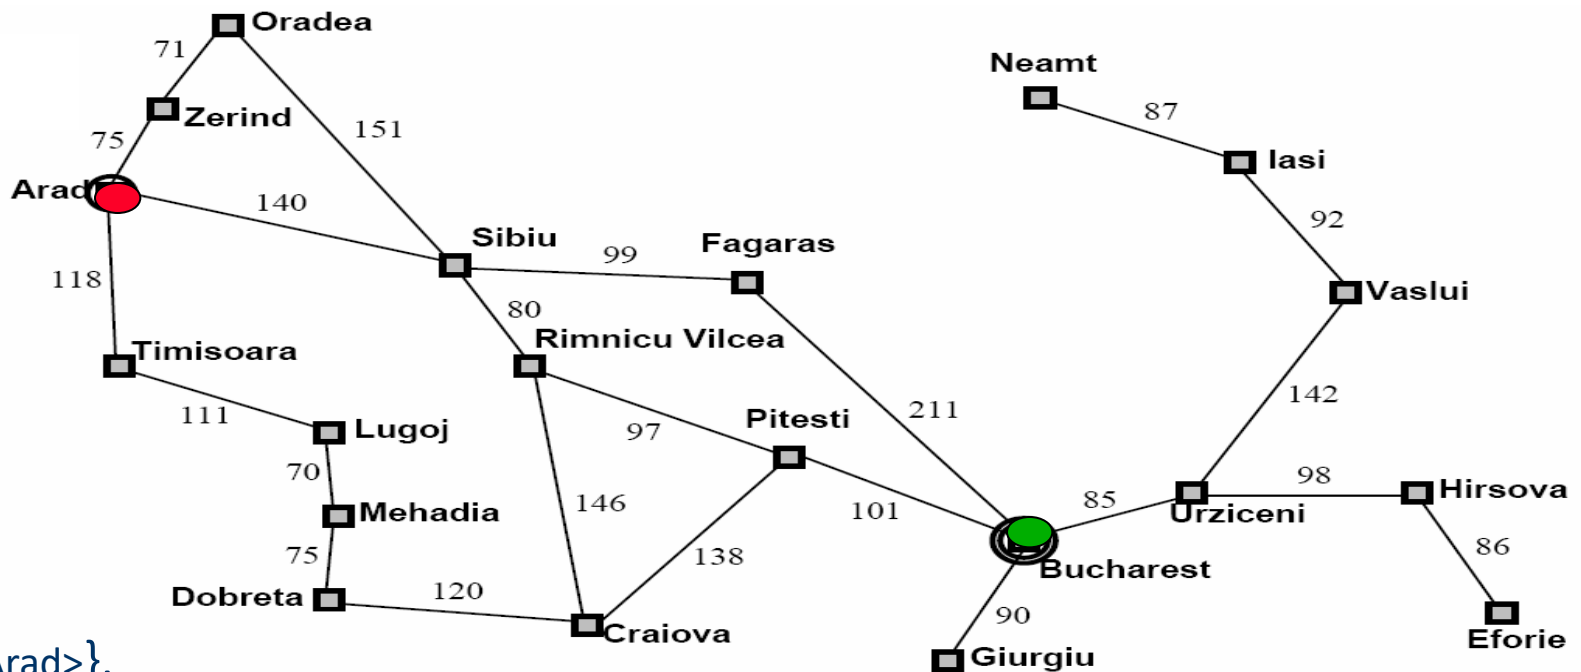
{<A,Z>, <A,T>, <A,S,A>, <A,S,O>, <A,S,F>, <A,S,R>}

{<A,Z>, <A,T>, <A,S,A>, <A,S,O>, <A,S,F>, <A,S,R,S>, <A,S,R,C>, <A,S,R,P>}

{<A,Z>, <A,T>, <A,S,A>, <A,S,O>, <A,S,F>, <A,S,R,S>, <A,S,R,C>, <A,S,R,P,R>, <A,S,R,P,C>, <A,S,R,P,B>}

**Solution:** Arad -> Sibiu -> Rimnicu Vilcea -> Pitesti -> Bucharest

**Cost:** 140 + 80 + 97 + 101 = 418



{<Arad>},

{<A,Z>, <A,T>, <A, S>},

{<A,Z>, <A,T>, <A,S,A>, <A,S,O>, <A,S,F>, <A,S,R>}

{<A,Z>, <A,T>, <A,S,O>, <A,S,F>, <A,S,R>, <A,S,A,Z>, <A,S,A,T>, <A,S,A,S>}

.....

**cycles** can cause non-termination!

... we deal with this issue later

# Selection Rule

---

The order paths are selected from OPEN has a critical effect on the operation of the search:

- Whether or not a solution is found
- The cost of the solution found.
- The time and space required by the search.

# How to select the next path from OPEN?

---

All search techniques keep OPEN as an ordered set and repeatedly execute:

- If OPEN is empty, terminate with failure.
  - Remove the **first** path from OPEN (OPEN is ordered!)
  - If the path leads to a goal state, terminate with success.
  - Extend the path (i.e. generate the successor states of the final state of the path) and put the new paths in OPEN.
- 
- The question of which path to select next from OPEN is now equivalent to the question of how do we order the paths on OPEN?

# Critical Properties of Search

---

- **Completeness**: will the search always find a solution (a path reaching a goal state) if a solution exists?
- **Optimality**: will the search always find the least cost solution? (when actions have costs)
- **Time complexity**: what is the maximum number of paths than can be expanded or generated?
- **Space complexity**: what is the maximum number of paths that have to be stored in memory?
  - Note that since each path is stored as its final state and a pointer to the previous state, the new paths  $\langle p, \text{succ} \rangle$  added to OPEN only require  $O(1)$  space, so we are only interested in the number of paths added to OPEN



# Uninformed Search Strategies

---

- These are strategies that adopt a fixed rule for selecting the next state to be expanded.
- The rule does not change, particular properties of the search problem being solved are ignored.
- Uninformed search techniques:
  - Breadth-First, Uniform-Cost, Depth-First, Depth-Limited, and Iterative-Deepening search

# Uninformed Search

---

- You would have seen breadth-first search and depth-first search in CSC263/265.
  - Graph nodes = state space states
  - Graph edges = state space actions
- In that course however, it is assumed that the graph we are searching is **explicitly represented** as an adjacency list (or adjacency matrix).
  - This won't work when there are an exponential number of nodes and edges.
- Similarly uniform cost search is like Dijkstra's algorithm, but without an explicitly represented graph.
- All of these algorithms are simple instantiations of our implicit graph search.

---

# Breadth-First Search

# Breadth-First Search

---

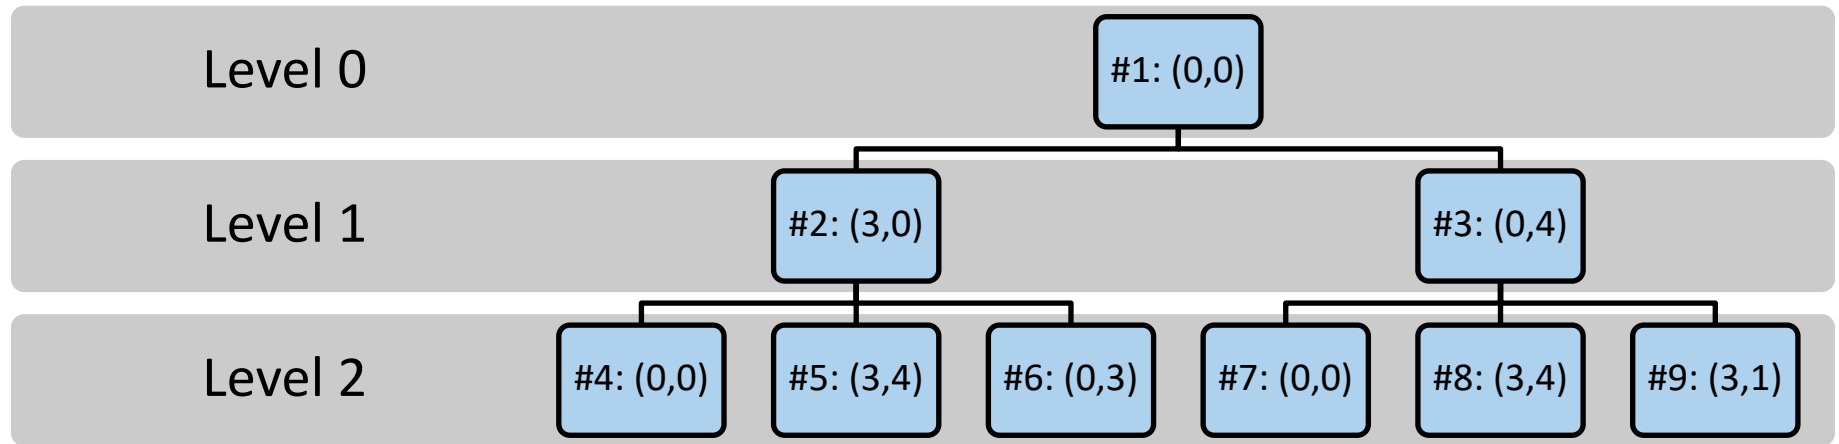
- Place the new paths that extend the current path at the **end** of OPEN. OPEN is a **queue**. Always extract first element of OPEN

WaterJugs. Start = (0,0), Goal = (\*,2)

**Red** = Expanded next. **Green** = newly added

1. OPEN = {<(0,0)>}
2. OPEN = {<(0,0),(3,0)>, <(0,0),(0,4)>}
3. OPEN = {<(0,0),(0,4)>, <(0,0),(3,0),(0,0)>, <(0,0),(3,0),(3,4)>, <(0,0),(3,0),(0,3)>}
4. OPEN = {<(0,0),(3,0),(0,0)>, <(0,0),(3,0),(3,4)>, <(0,0),(3,0),(0,3)>, <(0,0),(0,4),(0,0)>, <(0,0),(0,4),(3,4)>, <(0,0),(0,4),(3,1)>}

# Breadth-First Search



- Above we indicate only the state that each path terminates at. The path is from the root to that state.
- Breadth-First explores the search space level by level.

# Breadth-First Properties


---

## Completeness?

- The length of the path removed from OPEN is non-decreasing.
  - we replace each expanded path **p** of length  $k$  with a path that of length  $k+1$
  - **All** shorter paths are expanded prior to any longer path.
- If there is a solution, it is a path of length  $k$  for some  $k$ . Hence, we will eventually examine all paths of length  $< k$ , then all paths of length  $k$  and thus find a solution if one exists.
  - *What are we assuming here?*

## Optimality?

- By the above will find shortest length solution
  - least cost solution?
  - Not necessarily: shortest solution not always cheapest solution if actions have varying costs



**Breadth first Solution:** Arad -> Sibiu -> Fagaras -> Bucharest

**Cost:**  $140 + 99 + 211 = 450$

**Lowest cost Solution:** Arad -> Sibiu -> Rimnicu Vilcea -> Pitesti -> Bucharest

**Cost:**  $140 + 80 + 97 + 101 = 418$

# Breadth-First Properties

---

Measuring time and space complexity (**Note our search algorithm omits a simple optimization of Breadth-First Algorithm used in the recommended text. Hence our analysis is different.**)

- let  $b$  be the maximum number of successors of any state (**maximal branching factor**).
- let  $d$  be the depth of the shortest solution.
  - Path from root (at depth 0) to a state at level  $d$  has length  $d$

Time Complexity?

$$1 + b + b^2 + b^3 + \dots + b^{d-1} + b^d + b(b^d - 1) = O(b^{d+1})$$



# Breadth-First Properties

---

## Space Complexity?

- $O(b^{d+1})$ : If only the last path of length  $d$  satisfies  $d$ , then all paths of length  $d$  except the last one will be expanded by the search. Each such expansion will add up to  $b$  new paths to OPEN. So we can have up to  $b(b^d - 1)$  paths on open by the time we stop by expanding a path reaching a goal state

# Breadth-First Properties

---

Space complexity is a real problem.

- E.g., let  $b = 10$ , and say 100,000 paths can be expanded per second and each path requires 100 bytes of storage:

Depth	Paths	Time	Memory
1	1	0.01 millisec.	100 bytes
6	$10^6$	10 sec.	100 MB
8	$10^8$	17 min.	10 GB
9	$10^9$	3 hrs.	100 GB

- Typically run out of space before we run out of time in most applications.

---

# Uniform-Cost Search

# Uniform-Cost Search

---

- Keep OPEN ordered by increasing cost of the path.
- Always expand the least cost path.
- Identical to Breadth first if each action has the same cost.

# Uniform-Cost Properties

---

## Completeness?

- If each transition has costs  $\geq \epsilon > 0$ .
- The previous argument used for breadth first search holds: the cost of the path chosen to be expanded must be non-decreasing so eventually we will expand all paths with cost equal to the cost of a solution path.

## Optimality?

- Finds optimal solution if each transition has cost  $\geq \epsilon > 0$ .
- Explores paths in the search space in increasing order of cost. So must find minimum cost path to a goal before finding any higher costs paths leading to a goal

# Uniform-Cost Search. Proof of Optimality

---

Let us prove Optimality more formally.

# Uniform-Cost Search. Proof of Optimality

---

## Lemma 1.

Let  $c(p)$  be the cost of path  $p$  on OPEN (cost of the path). If  $p'$  is expanded after  $p$  then  $c(p) \leq c(p')$ .

Proof: First assume that  $p'$  is expanded **immediately after**  $p$ . There are 2 cases:

- a.  $p'$  was on OPEN when  $p$  was expanded:

*We must have  $c(p) \leq c(p')$  otherwise  $p'$  would have been selected for expansion rather than  $p$*

- b.  $p'$  was added to OPEN when  $p$  was expanded

*Now  $c(p) < c(p')$  since the path represented by  $p'$  extends the path represented by  $p$  and thus  $p'$  has cost at least  $\epsilon$  more than  $p$*

Now let  $p'$  be expanded **after**  $p$  (but not necessarily immediately after. This means we have expanded a sequence of paths  $p, p_1, p_2, \dots, p_k, p'$ . But we know that  $c(p) \leq c(p_1) \leq c(p_2) \dots \leq c(p_k) \leq c(p')$ . That is, we have that  $c(p) \leq c(p')$ .

# Uniform-Cost Search. Proof of Optimality

---

## Lemma 2.

When path  $p$  is expanded every path in the search space with cost strictly less than  $c(p)$  has already been expanded.

### Proof:

- Let  $p_k = \langle \text{Start}, s_1, \dots, s_k \rangle$  be any path with cost less than  $c(p)$ . We will prove that  $p_k$  has already been expanded (i.e., before  $p$  was extracted from OPEN to be expanded)
- Define a collection of paths  $p_i$  each of which is a prefix of the path  $p_k$ .  
 $p_0 = \langle \text{Start} \rangle,$   
 $p_1 = \langle \text{Start}, s_1 \rangle,$   
 $p_2 = \langle \text{Start}, s_1, s_2 \rangle,$   
 $\dots,$   
 $p_i = \langle \text{Start}, s_1, \dots, s_i \rangle,$   
 $\dots,$   
 $p_k = \langle \text{Start}, s_1, \dots, s_k \rangle.$ 
  - **Note that since each path is a prefix of the subsequent paths**  
 **$c(p_i) < c(p_{i+1}) < c(p_{i+2}) \dots < c(p_k)$**



# Uniform-Cost Search. Proof of Optimality

---

- Let  $p_i$  be the last path in this sequence of paths that has **already been expanded** by search.
  - If  $i = k$ , then  $p_k$  has **already been expanded** and we are done
  - If  $i < k$ , then we show that this is a contradiction, so this case cannot happen. So we are left with the case  $i=k$  and  $p_k$  must have already been expanded.
- 
- Now we prove that  $i < k$  is impossible (i.e., it yields a contradiction).
  - Since  $p_i$  is the last path on the sequence to already expanded,  $p_{i+1}$  has not yet been expanded. However,  $p_{i+1}$  is a successor of  $p_i$ , so it must have been added to OPEN when  $p_i$  was expanded. Hence,  $p_{i+1}$  must be on OPEN when path  $p$  was selected to be expanded, and thus  $c(p_{i+1}) \geq c(p)$  otherwise UCS (uniform cost search) would have expanded  $p_{i+1}$  instead of  $p$ .
  - On the other hand since  $i < k$ ,  $i+1 \leq k$  so  $c(p_{i+1}) \leq c(p_k) < c(p)$ .  
**Contradiction.** Hence the case  $i < k$  cannot happen, and  $p_k$  must **already have been expanded** when the Uniform cost search expands  $p$ .

# Uniform-Cost Search. Proof of Optimality

---

## Lemma 3.

The first time uniform-cost expands a path  $p$  whose final state is state  $S$ , it has found the minimal cost path to  $S$  (it might later find other paths to  $S$  but none of them can be cheaper).

## Proof:

- All cheaper paths have already been expanded, none of them terminated at  $S$ .
- All paths expanded after  $p$  will be at least as expensive, so no cheaper path to  $S$  can be found later.

So, when a path to a goal state is expanded the path must be optimal (lowest cost).

# Uniform Cost Search Visualizing the Proof

---

- Paths in the State Space ordered by cost

Path	Cost
<S>	0
<S, a>	1.0
<S,b>	1.0
<S,a,c>	1.5
<S,d>	2.0
...	...

# Uniform Cost Search Visualizing the Proof

---

- Uniform Cost Search expands paths in non-decreasing order of cost. So it can only go down this list.

LEMMA 1

Path	Cost
<S>	0
<S, a>	1.0
<S,a,b>	1.5
<S,a,b,c>	3.0
<S,a,b,d,e>	4.0
...	...

# Uniform Cost Search Visualizing the Proof

---

It does not miss any paths on this list. LEMMA 2

Path	Cost
<S>	0
<S, a>	1.0
<S,a,b>	1.5
<S,a,b,c>	3.0
<S,a,b,d>	4.0
...	...

- If <S,a,b,d> is expanded next, <S,a,b,c>, <S,a,b> ..., must all have already been expanded as all paths with lower cost than <S,a,b,d> must already be expanded

# Uniform Cost Search Visualizing the Proof

---

Thus working its way down such a list of paths the first path achieving the goal that UCS finds must be the cheapest path reaching the goal.

# Uniform-Cost Properties

---

## Time and Space Complexity?

- Uniform-Cost search (UCS) has complexity equal to the number of paths that have cost  $\leq C^*$  where  $C^*$  is the cost of an optimal solution.
- UCS has to expand all paths with cost  $< C^*$  and potentially all paths with cost  $= C^*$
- In the worst case we have  $O(b^{C^*/\epsilon})$  paths with cost  $\leq C^*$  so in the worst case the time and space complexity of UCS is  
$$O(b^{C^*/\epsilon})$$
- Why?  $C^*/\epsilon$  is the worst case number of states in a path of cost  $C^*$ , and each state might generate  $b$  alternate successor paths.

---

# Depth-First Search



# Depth-First Search

---

- Place the new paths that extend the current path at the **front** of OPEN. OPEN is a **stack**.

WaterJugs. Start = (0,0), Goal = (\*,2)

**Green** = Newly Added. **Red** expanded next

1. OPEN = {<(0,0)>}
2. OPEN = {<(0,0), (3,0)>, <(0,0), (0,4)>}
3. OPEN = {<(0,0), (3,0), (0,0)>, <(0,0), (3,0), (3,4)>, <(0,0), (3,0), (0,3)>, <(0,4), (0,0)>}
4. OPEN = {<(0,0), (3,0), (0,0), (3,0)>, <(0,0), (3,0), (0,0), (0,4)>, <(0,0), (3,0), (3,4)>, <(0,0), (3,0), (0,3)>, <(0,0), (0,4)>}

# Depth-First Search

Level 0

#1: (0,0)

Level 1

#2: (3,0)

(0,4)

Level 2

#3: (0,0)

(3,4)

(0,3)

Level 3

#4: (3, 0)

(0,4)

- Red paths are backtrack points (these paths remain on open).

# Depth-First Properties

---

## Completeness?

- Infinite paths? Cause incompleteness!
- Prune paths with cycles (duplicate states)

We get completeness **if state space is finite**

## Optimality?

No!

# Depth-First Properties

---

## Time Complexity?

- $O(b^m)$  where  $m$  is the length of the longest path in the state space.
- Very bad if  $m$  is much larger than  $d$  (shortest path to a goal state), but if there are many solution paths it can be much faster than breadth first.
- Can get lucky and bump into a solution quickly. Using heuristics to determine which successor to explore first can help in getting lucky.

# Depth-First Properties

---

- Depth-First Backtrack Points = unexplored siblings of paths along current path.
  - Each path can have at most  $b$  unexplored siblings.
  - These are the paths that remain on open after we extract a path to expand.

## Space Complexity?

- $O(bm)$ , linear space!
  - Only explore a single path at a time.
  - OPEN only contains the current path along with the **backtrack** points.
    - How many backtrack points are there for a depth  $k$  path?
- A significant advantage of DFS

---

# Depth-Limited Search

# Depth Limited Search

---

- Breadth first has space problems. Depth first can run off down a very long (or infinite) path.

## Depth limited search

- Perform depth first search but only to a pre-specified depth limit  $D$ .
  - THE ROOT is at DEPTH 0. ROOT is a path of length 0.
  - No path of length  $> D$  is placed on OPEN.
  - We “truncate” the search by looking only at paths of length  $D$  or less.
- Now infinite length paths are not a problem.
- But will only find a solution if a solution of  $\text{DEPTH} \leq D$  exists.

# Depth Limited Search

---


```
DL_Search(open, successors, goal?, maxd):
    open.insert(<start>)          #OPEN MUST BE A STACK FOR DFS
    cutoff = false
    while not open.empty():
        p = open.extract()       #remove path from OPEN
        if (goal?(p.final())):
            return (p,cutoff)    #p is solution
        if length(p) < maxd:     #Only successors if length(d) < maxd
            for succ in successors(p.final()):
                open.insert(<p,succ>)
        else:
            cutoff= true.        #some path was not
                                #expanded because of depth
                                #limit.

    return (null, cutoff)
```



# Depth Limited Search Example

Limit = 3



---

# Iterative Deepening Search

# Iterative Deepening Search

---

- Solve the problems of depth-first and breadth-first by extending depth limited search
- Starting at depth limit  $L = 0$ , we iteratively increase the depth limit, performing a depth limited search for each depth limit.
- Stop if a solution is found, or if the depth limited search failed without cutting off any paths because of the depth limit.
  - If no paths were cut off, the search examined all paths in the state space and found no solution  $\rightarrow$  no solution exists.

# Iterative Deeping Search

---

```
ID_Search(open, successors, goal?):  
    maxd = 0  
    while true:  
        (p, cutoff) = DL_Search(open, successors, goal?, maxd)  
        if p:  
            return p  
        elif not cutoff:           #no paths at deeper levels exit  
            return fail  
        else:  
            maxd = maxd + 1
```

# Iterative Deepening Search Example

---


Limit = 0



# Iterative Deepening Search Example


---

Limit = 1




# Iterative Deepening Search Example

Limit = 2



# Iterative Deepening Search Example

Limit = 3





# Iterative Deepening Search Properties

---

## Completeness?

- Yes. If the minimum length solution has length  $d$  then IDS will call DLS for each value of  $\text{maxd}$  in the range  $0—d$ .
- When it calls DLS with  $\text{maxd}=d$  then DLS will find and return a solution.

## Time Complexity?

# Iterative Deepening Search Properties

---

## Time Complexity

- $(d+1)b^0 + db^1 + (d-1)b^2 + \dots + b^d = O(b^d)$
- E.g.  $b=4, d=10$ 
  - $(11)*4^0 + 10*4^1 + 9*4^2 + \dots + 4^{10} = 1,864,131$
  - $4^{10} = 1,048,576$
  - Most paths terminate on on bottom layer.

# BFS can explore more states than IDS!

---

- For IDS, the time complexity is
  - $(d+1)b^0 + db^1 + (d-1)b^2 + \dots + b^d = O(b^d)$
- For BFS, the time complexity is
  - $1 + b + b^2 + b^3 + \dots + b^d + b(b^d - 1) = O(b^{d+1})$

E.g.  $b=4, d=10$

- For IDS
  - $(11)*4^0 + 10*4^1 + 9*4^2 + \dots + 4^{10} = 1,864,131$  (states generated)
- For BFS
  - $1 + 4 + 4^2 + \dots + 4^{10} + 4(4^{10} - 1) = 5,592,401$  (states generated)
  - In fact IDS can be more efficient than breadth first search: paths at limit are not expanded. BFS must expand all paths until it expands a goal state. So at the bottom layer it will add many path to OPEN before finding the path reaching a goal state.
  - With a simple optimization BFS can achieve the same time complexity as IDS, **but its time complexity is not better and its space complexity is much worse.**
  - In practice however BFS can be much better or much worse depending on the problem. In particular, more effective cycle checking can be employed with BFS. With IDS this more effective cycle checking will make the space complexity as bad as BFS—this removes the main advantage of IDS.

# Iterative Deepening Search Properties

---

## Space Complexity

- $O(bd)$  Still linear!

## Optimal?

- Will find shortest length solution which is optimal if costs are uniform.
- If costs are not uniform, we can use a “cost” bound instead.
  - Only expand paths of cost less than the cost bound.
  - Keep track of the minimum cost unexpanded path in each depth first iteration, increase the cost bound to this on the next iteration.
- This will need as many iterations of the search as there are distinct path costs. There can be more different path costs than there are different path lengths. So IDS using increasing costs can be significantly less efficient.

---

# Path/Cycle Checking

# Path Checking

---

If  $p_k$  is the path  $\langle s_0, s_1, \dots, s_k \rangle$  and we expand  $s_k$  to obtain child successor state  $c$ , we have

$$\langle s_0, s_1, \dots, s_k, c \rangle$$


As the path to “ $c$ ”.

We write such paths as  $\langle p, c \rangle$  where  $p$  is the prefix and  $c$  is the final state in the path.

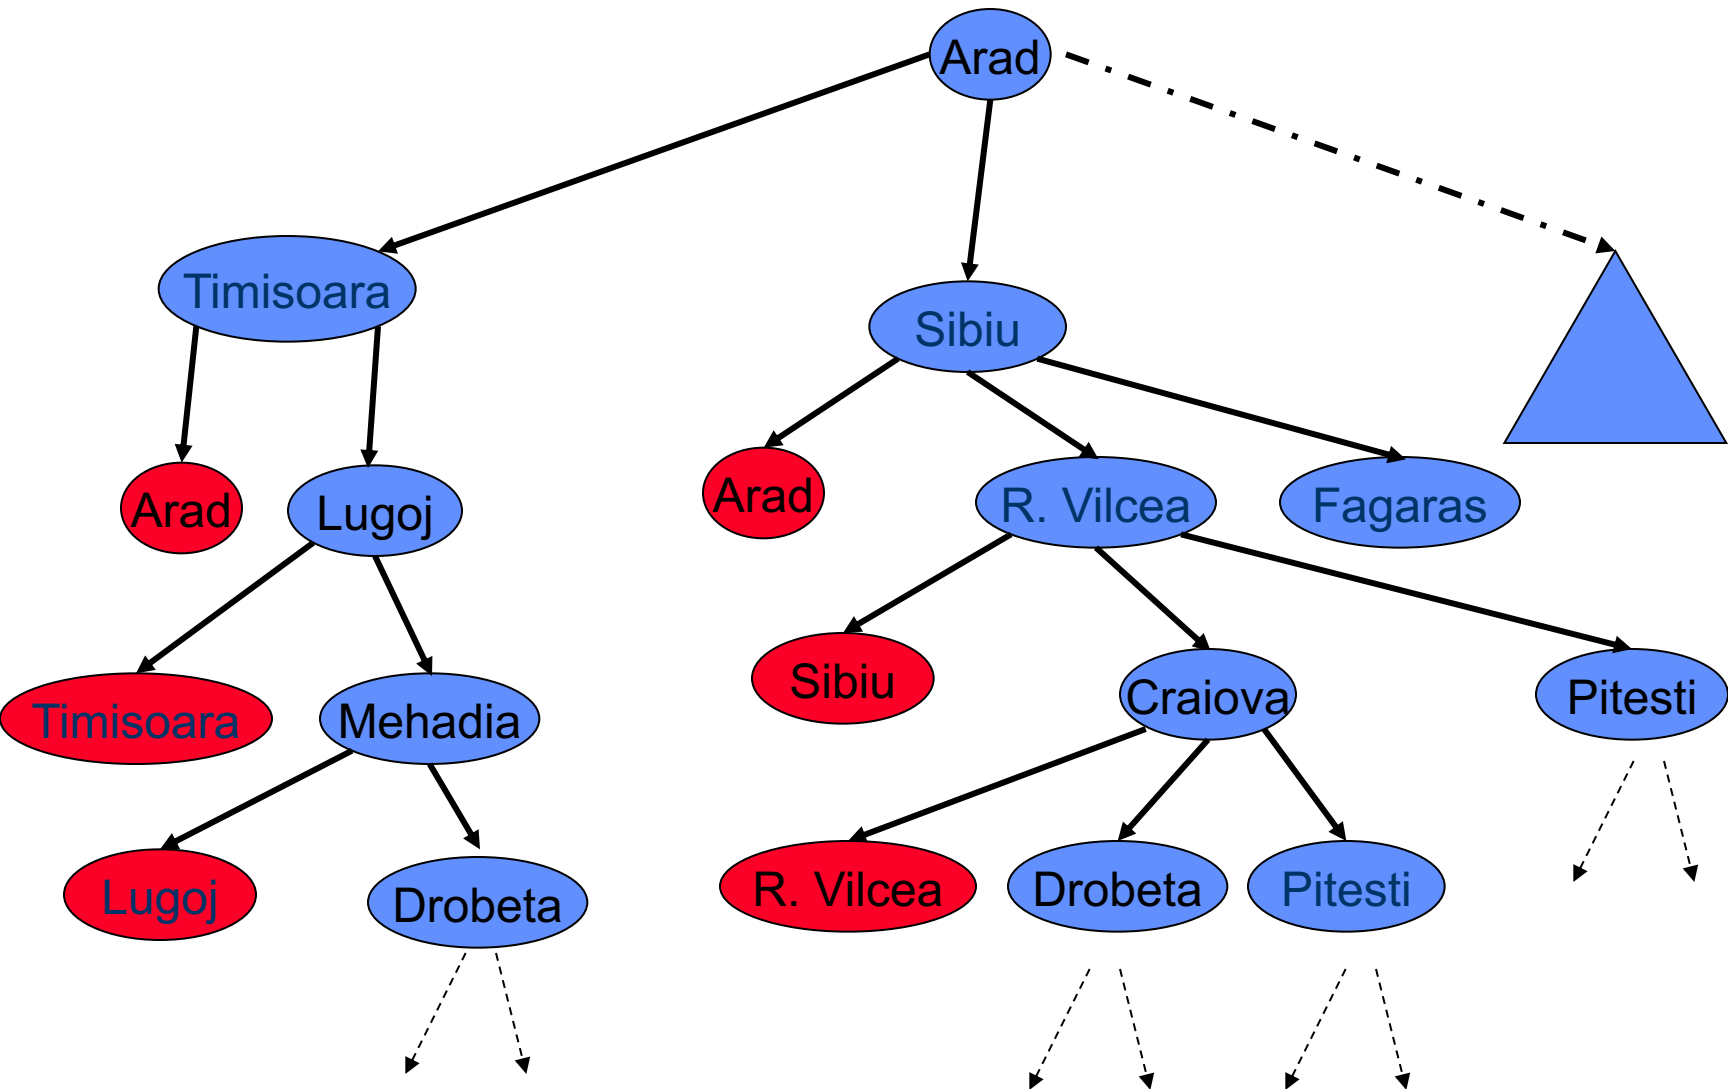
Path checking:

- Ensure that the state  $c$  is not equal to the state reached by any ancestor of  $c$  along this path.  $c \notin \{s_0, s_1, \dots, s_k\}$
- Paths are checked in isolation!

# Example: Arad to Neamt



# Path Checking Example





# Search with Path Checking

---

```
Search(open, successors, goal? ):  
    open.insert(<start>)  
    while not open.empty():  
        p = open.extract()           #remove path from OPEN  
        if (goal?(p.final())):  
            return p                 #p is solution  
        for succ in successors(state):  
            if not succ in <p>:      #Don't put cyclic paths on OPEN  
                open.insert(<p,succ>)  
    return false
```

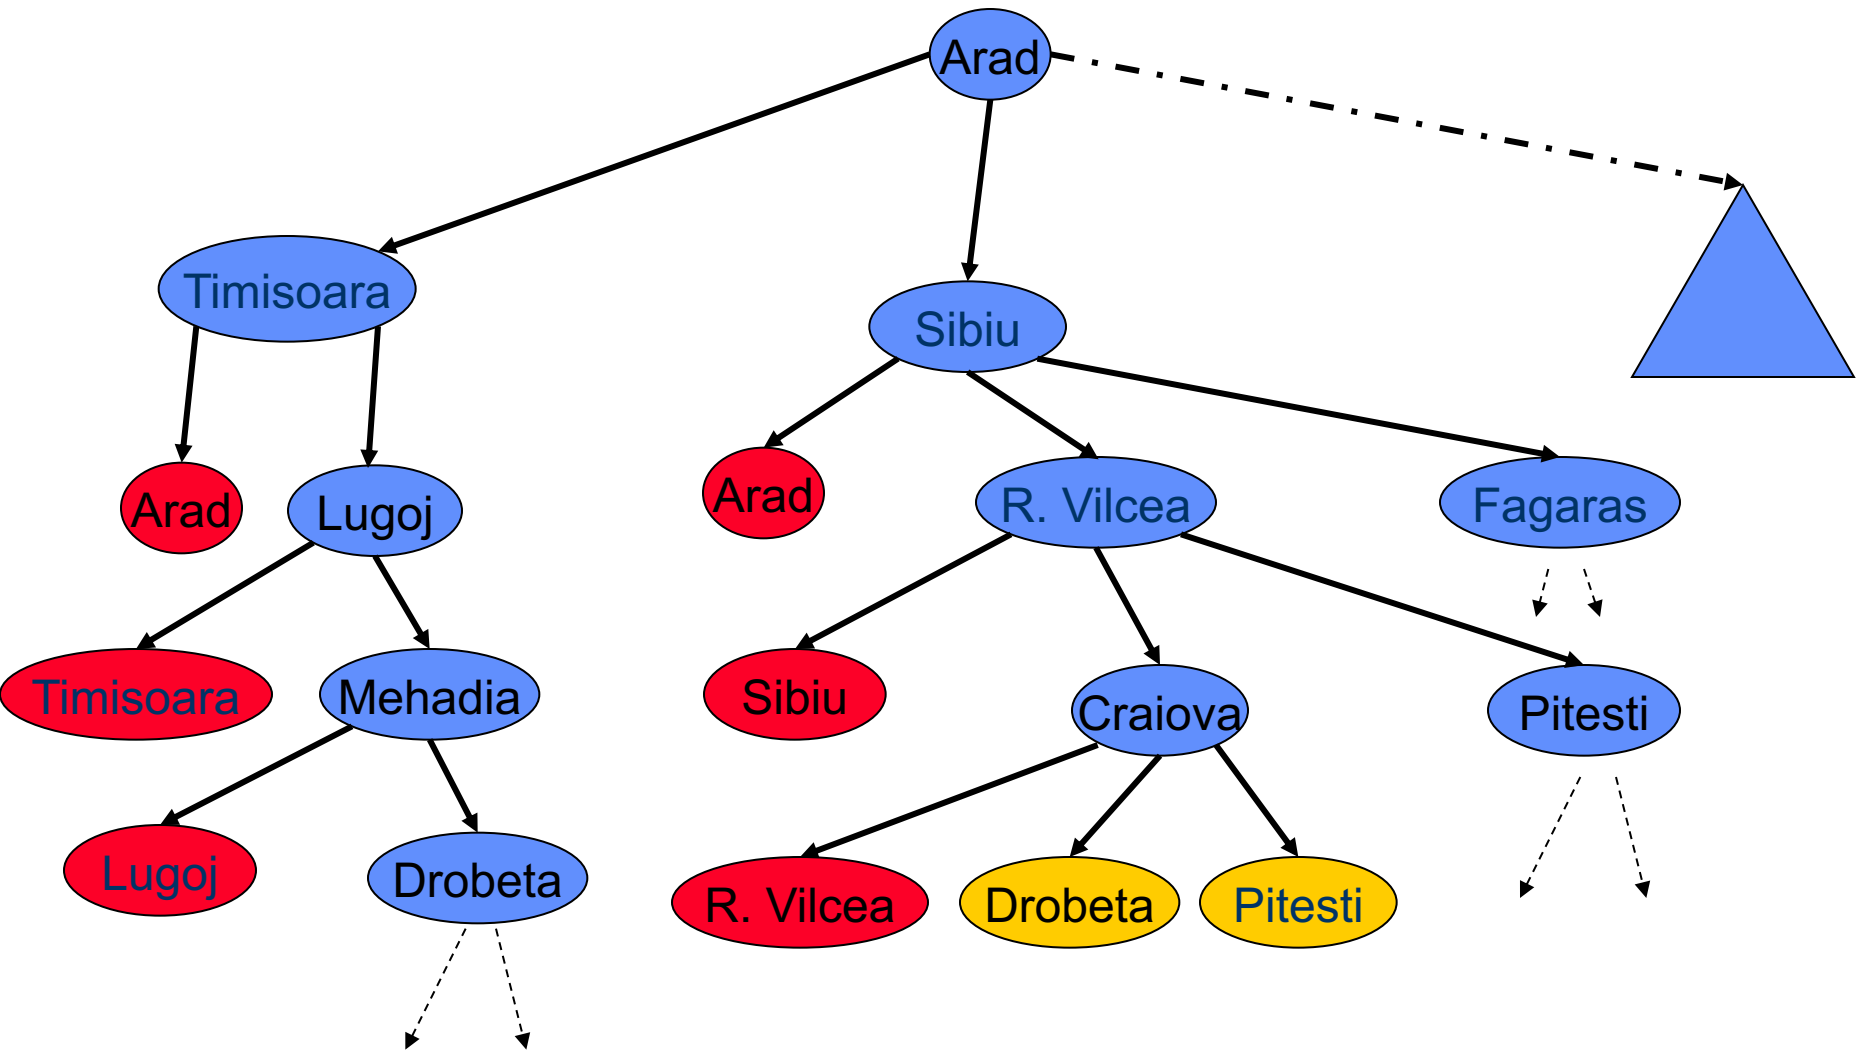
# Cycle Checking

---

## Cycle Checking

- Keep track of **all states** added to OPEN during the search (i.e., the final state of every path added to OPEN)
- When we expand  $p_k$  to obtain successor state  $c$ 
  - Ensure that  $c$  is not equal to **any** previously seen state.
  - If it is we do not add the path  $\langle p_k, c \rangle$  to OPEN.
- This is called **cycle checking**, or **multiple path checking**.
- What happens when we utilize this technique with depth-first search?
  - **What happens to space complexity?**

# Cycle Checking Example (BFS)



# Cycle Checking

---

- Higher space complexity (equal to the space complexity of breadth-first search). So good for BFS/UCS and other high space complexity searches. Bad for DFS/IDS as it destroys their linear space complexity.
- There is an additional issue when we are looking for an optimal solution
  - If we reject a path  $\langle p, c \rangle$  because we have previously seen state  $c$  via a different path it could be that  $\langle p, c \rangle$  is a cheaper path to  $c$  that we had previously seen.
  - Solution is to also keep track of the minimum cost path to each seen state.

# Cycle Checking

---

- Keep track of each state as well as the minimum known cost of a path to that state.
- If we find a longer path to a previously seen state, we don't add it to OPEN
- If we find a shorter path to a previously seen state, we add it to OPEN and
  - Remove other more expensive paths to the same state OR
  - **Lazily remove these more expensive paths if and only if we decide to expand them.** This will mean that when we select a path from OPEN it might be a redundant (more expensive) path, and we have to check for that.

# Cycle Checking—ensuring optimality

```
Search(open, successors, goal? ):  
  open.insert(<start>)  
  seen = {start : 0}           #seen is dictionary storing min cost  
  while not open.empty():  
    p = open.extract()  
    state = p.final()  
    if cost(p) <= seen[state]: #all p.final() on OPEN are in seen  
      #Expand only if p is cheapest known path to state  
      if (goal?(state)):  
        return p  
      for succ in successors(state):  
        if not succ in seen or cost(<p,succ>) < seen[succ]:  
          #Add this path to succ to open if this is the  
          #the first or the cheapest path to succ found so far  
          open.insert(<p,succ>)  
          seen[succ] = cost(<p,succ>)  
  return false
```

---

# Heuristic Search (Informed Search)

# Heuristic Search

---

- In **uninformed search**, we don't try to evaluate which of the paths on OPEN are most promising. We never “look-ahead” to the goal.

E.g., in uniform cost search we always expand the cheapest path. We don't try to estimate the cost of extending that path to reach a goal state.

- Often we have some other knowledge about the merit of paths, e.g., going the wrong direction in Romania, and thus can have an estimate of the cost of extending the path to reach a goal state.



# Heuristic Search

---

Merit of an OPEN path: different notions of merit.

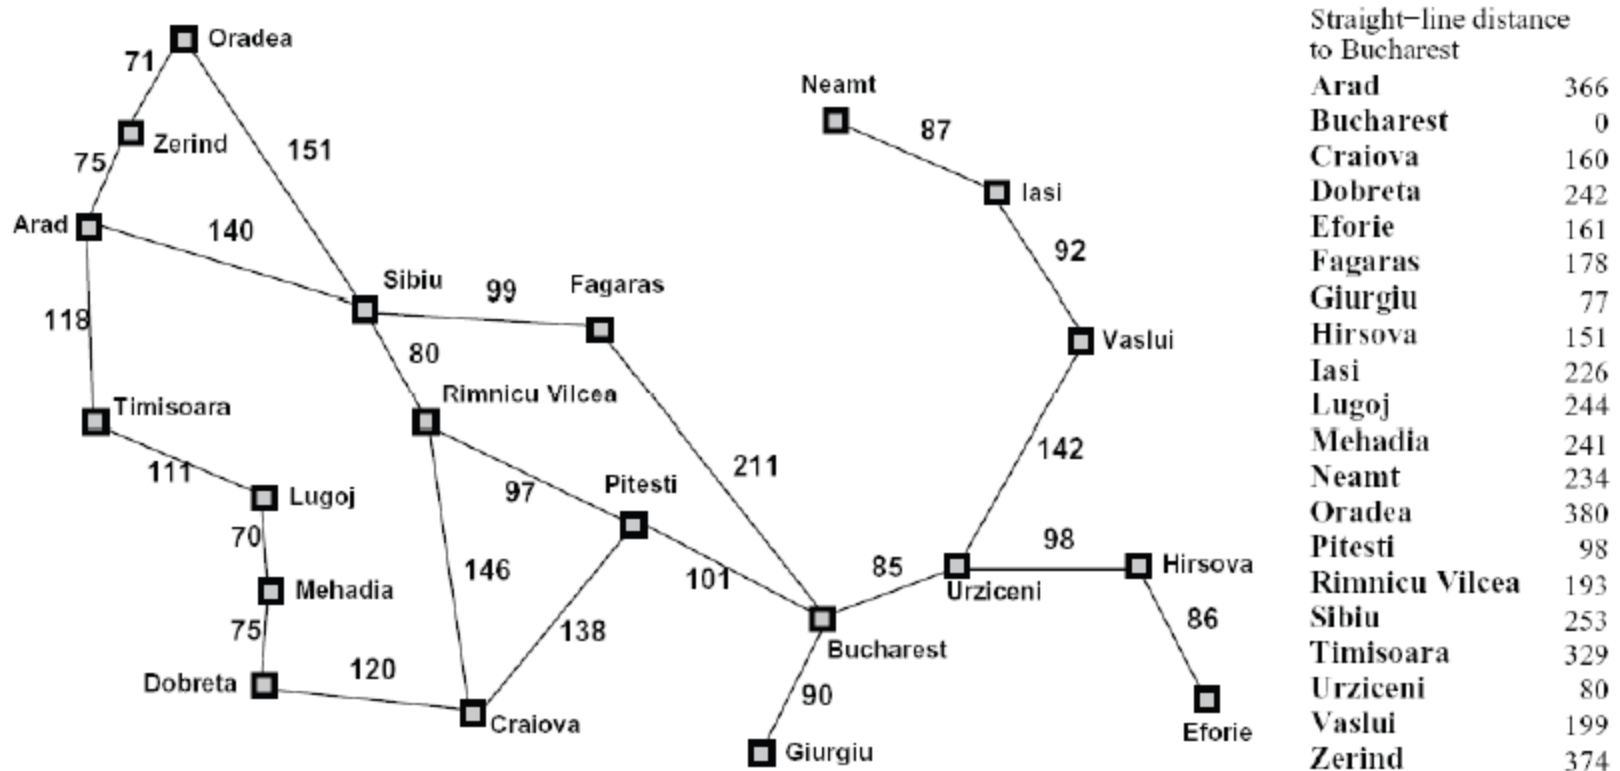
- If we are concerned about the **cost of the solution**, we might want to consider how costly it is to extend that path so that it reaches a goal state (i.e., the cost of getting from the path's final state to a goal).
- If we are concerned about **minimizing computation** in search we might want to consider how quickly the search will be able to extend that path so that it reaches a goal state (i.e., how hard is it to get from the path's final state to a goal).
- We will focus on the “**cost of solution**” notion of merit.

# Heuristic Search

---

- The idea is to develop a domain specific heuristic function  $h(p)$ .
- $h(p)$  **guesses** the cost of getting to the goal from the final state of a path  $p$ .
- $h(p)$  is a function only of  $p.\text{final}()$ !  
If  $p1.\text{final}() = p2.\text{final}()$  then  $h(p1) = h(p2)$ 
  - Hence we will often specify the function  $h$  by specifying its value on each state rather than on different paths.
- There are different ways of guessing this cost in different domains.
  - heuristics are **domain specific**.
  - Alpha Go exploited machine learning techniques to compute the heuristic estimate of a state (go board configuration)

# Example: Euclidean distance



Planning a path from Arad to Bucharest, we can utilize the **straight line distance from each city to our goal**. This lets us plan our trip by picking cities at each time point that minimize the distance to our goal.

# Heuristic Search

---

- If  $h(p_1) < h(p_2)$  this means that we **guess** that it is cheaper to get to the goal from  $p_1.\text{final}()$  than from  $p_2.\text{final}()$
- We require that
  - $h(p) = 0$  for every path  $p$  whose final state satisfies the goal.  
Zero cost of extending a path to reach a goal state when that path already reaches a goal state.


# Using only $h(p)$ : Greedy best-first search

- We use  $h(p)$  to rank the paths on OPEN
  - Always expand path with lowest  $h$ -value.
- We are greedily trying to achieve a low cost solution.
- However, this method **ignores the cost of  $p$** , so it can be lead astray exploring paths that cost a lot but seem to be close to the goal:

→ step cost = 10

→ step cost = 100


$h(n1) = 20$



$h(n3) = 10$

Note we write  $h(n3)$  instead of  $h(<S, n3>)$ !

# Greedy best-first search example



Straight-line distance to Bucharest	
Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	178
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

# Greedy best-first search

---

- Greedy search can be very efficient in practice at finding solutions...require developing a good heuristic.
- Greedy search is incomplete—if it fails to find a solution this does not mean a solution does not exist.
- The solution returned by a greedy search can be very far from optimal—not easy to control greedy search so that it finds close to optimal solutions.

One method is to add some randomness to the heuristic and compute many different solutions taking the best.

# A\* search


---

- Take into account the cost of the path as well as the heuristic estimate of the cost of extending the path to reach a goal state.
- Define an evaluation function  $f(n)$  (we use  $f(n)$  instead of  $f(p)$  as this is more traditional,  $n$  is a **path**)  
 $f(n) = g(n) + h(n)$ 
  - $g(n)$  is the cost of the path  $n$ .
  - $h(n)$  is the heuristic estimate of the cost of achieving the goal from  $n$ .final().





# A\* search

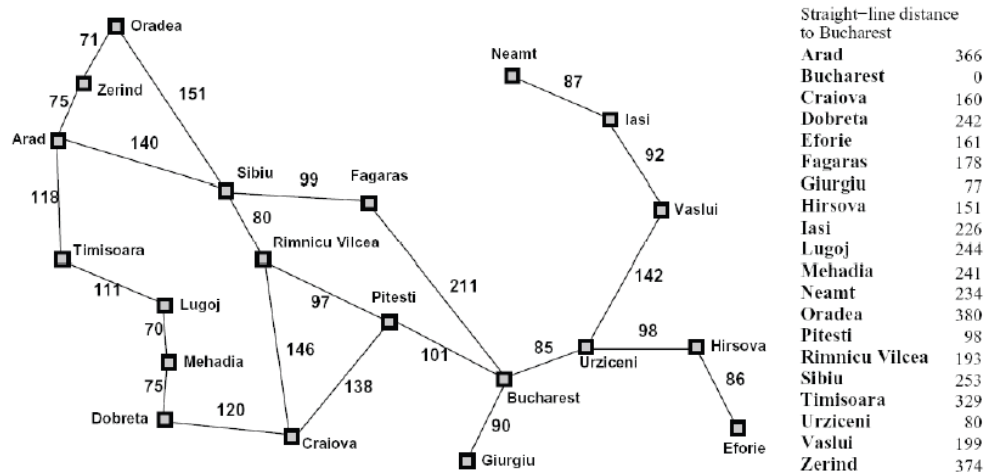
- $h(n)$  is a function only of  $n.\text{final}()$ —the end state reached by path  $n$ . So if  $n1.\text{final}() = n2.\text{final}()$  then  $h(n1) = h(n2)$
- $g(n)$  is the sum of the costs of the actions on the path  $n$ .
- Always expand the path with lowest  $f$ -value on OPEN.
- The  $f$ -value,  $f(n)$  is an estimate of the cost of getting to the goal via the path  $n$ .
  - I.e., we first follow the path  $n$  then we try to get to the goal.  $f(n)$  estimates the total cost of such a solution.




# A\* example



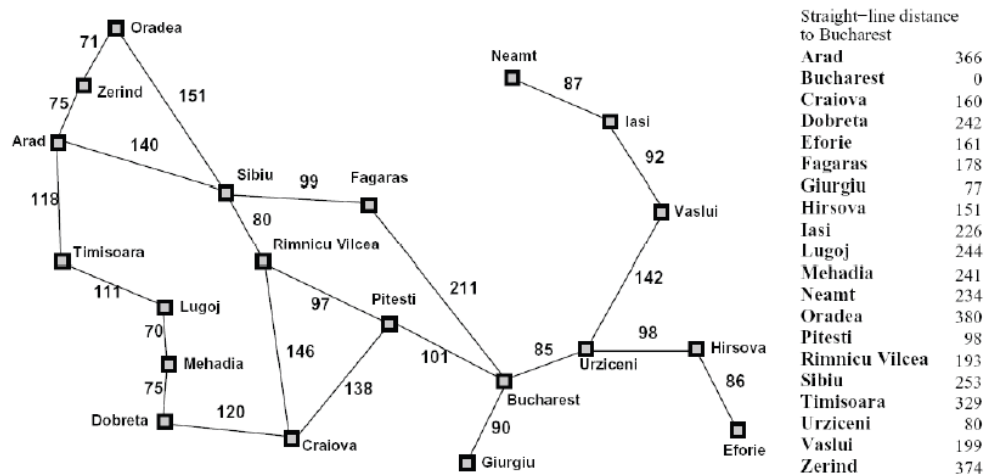
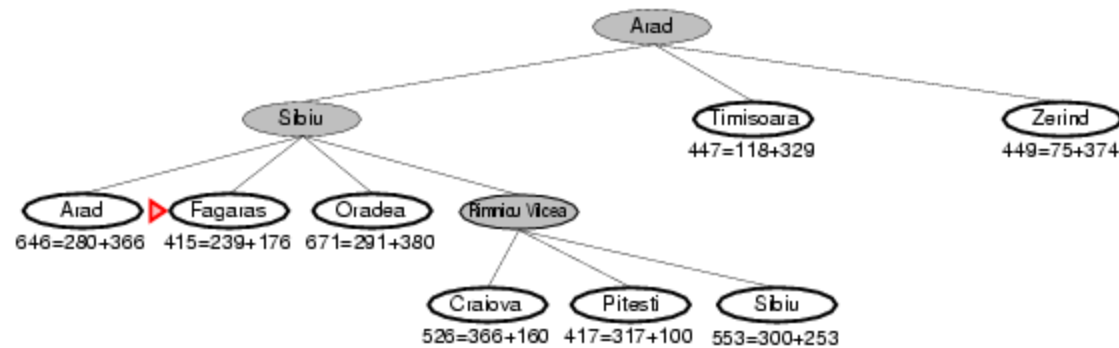
# A\* example



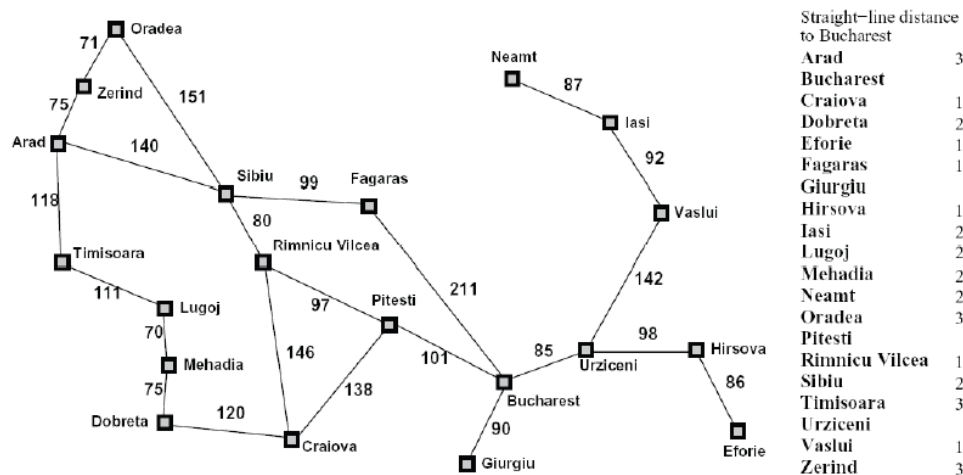
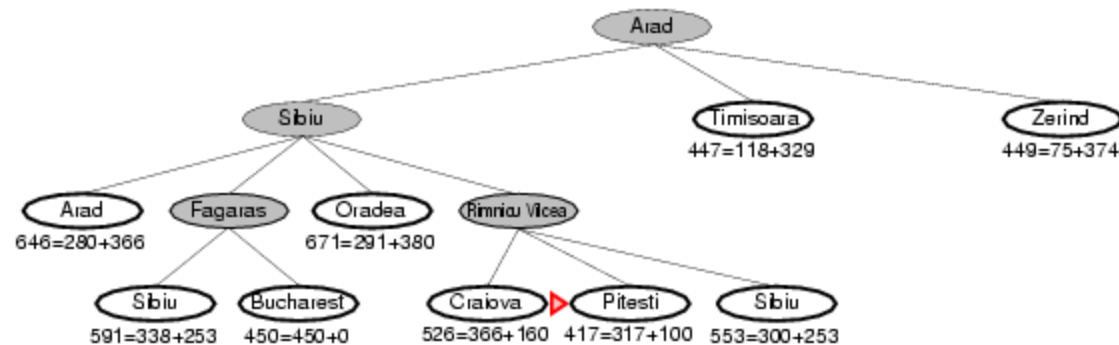
# A\* example



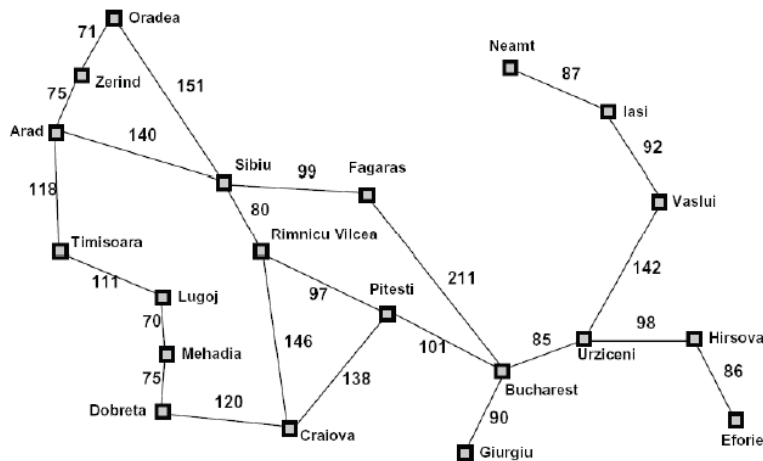

# A\* example



# A\* example



# A\* example



Straight-line distance  
to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	178
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

# A\* search Breaking Ties

---

- Suppose we have multiple paths with the lowest f-value. Is it important which of these we choose?
- **YES!** When we have admissible heuristics (the notion of admissibility is coming next) this can give an exponential speed-up.



# A\* search Breaking Ties

---

- Given that  $h(n)$  is an estimate of “how far away the goal is from  $n.final()$ ”, and  $g(n)$  is the cost already incurred for traversing the path  $n$ .
  - It makes intuitive sense to break ties by largest  $g$ -value (or equivalently by smallest  $h$ -value)
  - That is, if  $f(n1) = f(n2)$ 
$$f(n1) = g(n1) + h(n1) = g(n2) + h(n2) = f(n2)$$
so  $g(n1) > g(n2) \Leftrightarrow h(n1) < h(n2)$
  - So intuitively we would prefer to expand  $n1$  as it is closest to the goal, and is estimated to have equal total cost.
- Breaking ties in this way can reduce A\*'s search by an exponential factor! (more on this in tutorial).

# Formal Properties of A\*

---

- We want to analyze the behavior of the resultant search.
  - Completeness, time and space, optimality?
- We start with a simple result concerning paths in the search space that will be useful in proving other results.

# Reminder—Paths in the search space

---

- Remember that a path is a sequence of states  $\langle s_0, \dots, s_k \rangle$  generated by a sequence of actions  $\langle a_0, \dots, a_{k-1} \rangle$ . Where  $s_{i+1}$  is the result of applying action  $a_i$  to state  $s_i$ .
- The cost of the path is the sum of the costs of the actions  $a_i$ .
- In our discussion we will mostly leave the sequence of actions implicit—it exists and it determines the cost of the path, but we only need to talk about the states.

# Paths in the Search Space.

---

**Proposition 1.** For every path  $n = \langle s_0, \dots, s_k \rangle$  in the search space, at any time in the running of  $A^*$  either:

- a.  $\langle s_0, \dots, s_k \rangle$  has been expanded by  $A^*$  OR
- b. Some prefix  $\langle s_0, \dots, s_i \rangle$  of  $n$  is on OPEN.

**Proof:** By induction over the number of path expansions  $A^*$  has done.

**Base Case** (zero path expansions):  $A^*$  always starts off with  $\langle s_0 \rangle$  on OPEN (where  $s_0$  is the initial state). Every path  $n = \langle s_0, \dots, s_k \rangle$  has  $\langle s_0 \rangle$  as a prefix so (b) is true.

# Paths in the Search Space.

---

**Induction.** Assume true after  $k$  path expansions, prove true after the  $k+1$  expansion.

By induction after the  $k$ -th expansion either (a) or (b) is true. If (a) then  $n$  has already been expanded, if (b) then a prefix of  $n$  is on open after the  $k$ -th expansion.

If (a) is true ( $n$  has been expanded) it clearly stays expanded after  $k+1$  path expansions.

If (b) is true then let  $n_i = \langle s_0, \dots, s_i \rangle$  be the prefix of  $n$  that is on OPEN after  $k$  path expansions. Two possibilities:

1.  $A^*$  expands some path  $x \neq n_i$  as the  $k+1$  expansion. In this case  $n_i$  a prefix of  $n$  is still on OPEN.
2.  $A^*$  expands  $n_i$  as the  $k+1$  path expanded. Then either
  - a.  $n_i = n$  and now (a) is true— $n$  has been expanded OR
  - b. One of the successors of  $n_i$  is  $n_{i+1} = \langle s_0, \dots, s_{i+1} \rangle$ . So now  $n_{i+1}$  is on OPEN, and again (b) is true, i.e., a prefix of  $n$  remains on OPEN.

## Some Properties of $A^*$ depend on $h(n)$

---

- We want to analyze the behavior of the resultant search.
  - Completeness, time and space, optimality?
- We can obtain some results, but others depend on the properties of  $h(n)$ .
- First, what happens if there is no path from the initial state to a state satisfying the goal (i.e., there is no solution)?

# No Solution

---

- A\* as well as all of the uninformed search algorithms, depth first (DFS), breadth first (BFS), uniform cost (UCS), and iterative deepening (IDS) have the same behaviour when there is no solution:
  - If there are an **infinite** number of different states reachable from the initial state then these algorithms never terminate.
  - If there are a finite number of different reachable states **and** we do cycle checking (either path checking or full cycle checking) then we will eventually terminate correctly declaring that there is no solution. (Note costs are always  $\geq \epsilon > 0$ )

# Solutions

---

- Now consider the case where there are solutions
- This case has been covered when we analyzed each search algorithm to determine whether or not the algorithm is complete.
- Complete: BFS, UCS, IDS (and as we will prove next  $A^*$ )
  - Cycle checking not required for completeness!
- Incomplete DFS, DLS.
  - DFS—with a finite number of states reachable from the initial state and path-checking is complete.



# Completeness of A\*

---

**Theorem 1.** A\* will always find a solution if one exists (i.e., it is complete) as long as

- Every state has a finite number of successors (the branching factor is finite)—**note state space might still be infinite.**
- Every action has finite cost  $\geq \epsilon > 0$ .
- $h(n)$  is finite for every path  $n$  that can be extended to reach a goal state.

**Proof:** If a solution path  $n$  exists, then by **Proposition 1** at all times either

(a)  $n$  been expanded by A\* OR

(b) a prefix  $n$  is on OPEN

If (a) then A\* has halted and returned  $n$  as the solution and the theorem holds.

# Completeness of $A^*$

---

If (b) then let the prefix on OPEN be  $n_i$   
 $n_i$  must have a finite f-value. It has a finite cost (g-value) and its h-value is finite.

As  $A^*$  continues to run, the f-value of the paths on OPEN eventually increase.

At every iteration we remove a path of lowest f-value and replace it with its successors that have larger g-values. So the g-values of the paths on OPEN must eventually increase without bound as  $A^*$  continues to run, and thus so must the f-values.

**Note.** The h-value of the successors might be lower, so the new successors might have lower f-value. But as we continue eventually the g-value increases so much that the f-values of all paths on OPEN must increase.

# Completeness of $A^*$

---

So, eventually either

- (a)  $A^*$  terminates because it found a solution OR
- (b)  $n_i$  becomes the path on OPEN with lowest f-value

If (a) then the theorem holds— $A^*$  finds a solution.

If (b) then  $n_i$  is expanded and either

- 1)  $n_i = \mathbf{n}$  and  $A^*$  returns  $\mathbf{n}$  as a solution and the theorem holds OR
- 2)  $n_i$  is replaced by its successors one of which is a longer prefix of  $\mathbf{n}$ ,  $n_{i+1}$

## Completeness of $A^*$

---

Applying the same argument to  $n_{i+1}$  we see that if  $A^*$  continues to run without finding a solution it will eventually expand every prefix of the path  $\mathbf{n}$  and it must terminate when it expands  $\mathbf{n}$ .

# Completeness of $A^*$

---

**Note.**  $A^*$  will eventually find a solution even in infinite search spaces (as long as the branching factor is finite), even if we don't use cycle checking, and no matter what function  $h$  we use (as long as  $h$  is finite on any path that can be extended to reach a goal state)

But we want to go beyond completeness and develop conditions that ensure that  $A^*$  finds **optimal solutions**.

One condition that ensures optimality is when the heuristic  $h$  is **admissible**.

# Conditions on $h(n)$ : Admissible

---

- We assume that the cost of any action  $a$  is  $\geq \epsilon > 0$ : the cost of any transition is greater than zero and can't be arbitrarily small.

Let  $h^*(n)$  be the **cost of an optimal path** from  $n$ .final() to a goal state. Then an **admissible** heuristic satisfies the condition

$$h(n) \leq h^*(n)$$

- An admissible heuristic **under-estimates** the true cost to reach the goal, i.e., it is **optimistic**
- Hence  $h(g) = 0$ , for any path  $g$  reaching a goal state  $g$  since  $h^*(g) = 0$
- Also define  $h^*(n) = \infty$  if there is no path from  $n$  to a goal state (note that  $h(n)$  need not be infinite)

# Intuition behind admissibility

---

$h(n) \leq h^*(n)$  means that the search won't miss any promising paths.

- If it really is cheap to get to a goal via  $n$  (i.e., both  $g(n)$  and  $h^*(n)$  are low), then  $f(n) = g(n) + h(n)$  will also be low, and although the search might ignore  $n$  in favour of paths with lower  $f$ -value, eventually it will run out of such paths and return to expanding  $n$ .
- This can be formalized to show that admissibility implies optimality.

# Admissible Heuristics → Optimal solutions

---

- Now we prove formally that if  $h(n)$  is admissible, then  $A^*$  (ordering OPEN by  $f(n) = g(n) + h(n)$ ) always finds an optimal solution.
- As before we must assume that
  - No state has an infinite number of successors, i.e., the branching factor  $b$  is finite. (But the state space can be infinite)
  - The cost of any action is  $\geq \epsilon > 0$ .



# Admissible Heuristics → Optimal solutions

---

- If there exists a solution, then there must exist an optimal solution (perhaps many).
- Start with a few observations about any optimal path.
- Let  $\mathbf{n} = \langle s_0, s_1, \dots, s_n \rangle$  be an optimal solution, i.e., a path from the initial state  $s_0$  to a state  $s_n$  that satisfies the goal such that  $\mathbf{n}$  has cost  $\leq$  cost of any other path to a goal satisfying state.

- Let

$$\mathbf{n}_0 = \langle s_0 \rangle$$

$$\mathbf{n}_1 = \langle s_0, s_1 \rangle$$

...

$$\mathbf{n}_i = \langle s_0, s_1, \dots, s_i \rangle$$

...

$$\mathbf{n}_k = \langle s_0, s_1, \dots, s_k \rangle$$

That is, the prefixes of various lengths of this optimal path

# Admissible Heuristics → Optimal solutions

---

- Observe that
  - $n_i = \langle s_0, s_1, \dots, s_i \rangle$  is an optimal path from state  $s_0$  to state  $s_i$
  - $\langle s_{i+1}, \dots, s_n \rangle$  is an optimal path from state  $s_{i+1}$  to a goal state.  
That is, there is no cheaper way of getting from  $s_{i+1}$  to a goal state than  $\langle s_{i+1}, \dots, s_n \rangle$
  - Every sub-path  $\langle s_i, \dots, s_j \rangle$  is an optimal path from  $s_i$  to  $s_j$ .

If not we could replace the sub-path with a cheaper path and obtain a cheaper path from  $s_0$  to a goal state which would contradict our assumption that  $\mathbf{n} = \langle s_0, s_1, \dots, s_n \rangle$  is optimal.

# Admissible Heuristics → Optimal solutions

---

**Proposition 2:**  $A^*$  with an **admissible heuristic** never expands a path with f-value greater than the cost of an optimal solution.

**Proof.** Let  $C^*$  be the cost of an optimal solution.

We know that  $A^*$  always expands a path on OPEN that has lowest f-value.

Let  $n = \langle s_0, s_1, \dots, s_n \rangle$  be an optimal solution (so  $s_n$  satisfies the goal condition).

So  $\text{cost}(n) = \text{cost}(\langle s_0, s_1, \dots, s_n \rangle) = C^*$

By **Proposition 1**, we know that since  $A^*$  terminates if it expands  $n$ , at every stage OPEN always contains a prefix of  $n = \langle s_0, s_1, \dots, s_n \rangle$ , i.e., one of  $n_0 \dots n_k \dots n$

So  $A^*$  never expands a path with f-value greater than the f-value of that prefix.

Now we observe that when the heuristic is admissible for any prefix of  $n$  say  $n_i$  we have

$$C^* \geq f(n_i)$$

So if this is true then  $A^*$  never expands a path with f-value  $> f(n_i)$ , i.e., with f-value  $> C^*$  and the proposition is proved. Now we prove that  $C^* \geq f(n_i)$

# Admissible Heuristics → Optimal solutions

---

$$\begin{aligned}C^* &= \text{cost}(\langle s_0, s_1, \dots, s_n \rangle) \\&= \text{cost}(\langle s_0, \dots, s_i \rangle) + \text{cost}(\langle s_i, \dots, s_n \rangle) \\&= g(n_i) + h^*(n_i) & (1) \\&\geq g(n_i) + h(n_i) = f(n_i) & (2)\end{aligned}$$

(1)  $g(n_i)$  is the  $\text{cost}(n_i) = \text{cost}(\langle s_i, \dots, s_i \rangle)$

$h^*(n_i)$  is the cost of an optimal path from  $s_i$  to any goal state, which must be equal to  $\text{cost}(\langle s_i, \dots, s_n \rangle)$  since  $\langle s_0, \dots, s_n \rangle$  is optimal

(2)  $h^*(n_i) \geq h(n_i)$  since  $h$  is admissible.

# Admissible Heuristics → Optimal solutions

---

**Theorem 2.**  $A^*$  with an admissible heuristic always find an optimal solution if a solution exists.

**Proof:**

- If a solution exists then by **Theorem 1**,  $A^*$  will terminate by expanding some solution path  $n$ .
- By **Proposition 2**,  $f(n) \leq C^*$  (the cost of an optimal solution) since  $A^*$  never expands a path with cost  $> C^*$
- Since  $n$  is a solution  $h(n) = 0$  so  $f(n) = g(n) = \text{cost}(n)$ .
- We also have that  $C^* \leq \text{cost}(n) = f(n)$  (since no solution can have lower cost than the optimal).
- So  $\text{cost}(n) = C^*$ , i.e.,  $A^*$  returns an optimal solution.

# Consistency (aka monotonicity)

---

- A stronger condition than  $h(n) \leq h^*(n)$ .
- A **monotone/consistent** heuristic **h** satisfies the triangle inequality: for all paths  $n1, n2$  and for all actions we have that

$$h(n1) \leq C(n1.final(), a, n2.final()) + h(n2)$$

Where  $C(n1.final(), a, n2.final())$  means the cost of getting from the final state of  $n1$  to the final state of  $n2$  via action  $a$ .

- Note that there might be more than one transition (action) between  $n1$  and  $n2$ , the inequality must hold for all of them.
- Monotonicity implies admissibility.
  - (forall  $n1, n2, a$ )  $h(n1) \leq C(n1, a, n2) + h(n2) \rightarrow$  (forall  $n$ )  $h(n) \leq h^*(n)$

# Consistency → Admissible

---

- **Assume consistency:**  $h(n) \leq c(n,a,n_2) + h(n_2)$

**Prove admissible:**  $h(n) \leq h^*(n)$

**Proof:**

If no path exists from  $n$  to a goal then  $h^*(n) = \infty$  and  $h(n) \leq h^*(n)$

Else let  $\langle s_n, s_{n+1}, \dots, g \rangle$  be an OPTIMAL path from  $n$  to a goal state  $g$ .

Note the cost of this path is  $h^*(n)$ , and each subpath  $n_i = \langle s_{n+i}, \dots, g \rangle$  has cost equal to  $h^*(n_i)$ .

Prove  $h(n) \leq h^*(n)$  by induction on the length of this optimal path.

**Base Case:  $n = g$**

By our conditions on  $h$ ,  $h(n)$  and  $h(n)^*$  are equal to zero so

$$h(n) \leq h(n)^*$$

**Induction Hypothesis:  $h(n_1) \leq h^*(n_1)$**

$$h(n) \leq c(n, a_1, n_1) + h(n_1) \leq c(n, a_1, n_1) + h^*(n_1) = h^*(n)$$

# Consequences of monotonicity

---

1. The sequence of  $f$ -values of the paths expanded by  $A^*$  is non-decreasing. That is, if  $n_2$  is expanded **after**  $n_1$  by  $A^*$  then  $f(n_1) \leq f(n_2)$ . (Not necessarily true for an admissible heuristic as that heuristic might be non-monotone)
2. With a monotone heuristic, the first time  $A^*$  expands a path  $\mathbf{n} = \langle s_0, \dots, s_n \rangle$  that reaches the state  $s_n$ ,  $\mathbf{n}$  must be a minimum cost path to  $s_n$ 
  - This means that cycle checking need not keep track of the minimum cost of getting to a state found so far. It can simply prune all future paths to a state if that state has already been reached by an expanded path.
  - Again this is not necessarily true for an admissible heuristic. So with admissible heuristics that are not monotone (or not known to be monotone) we have to keep track of the cost of getting to states as well as the states visited during cycle checking.



## Complete.

- Yes, as we saw in **Theorem 1**.

## Time and Space complexity.

- When  $h(n) = 0$ , for all  $n$   $h$  is monotone.
  - A\* becomes uniform-cost search!
- It can be shown that when  $h(n) > 0$  for some  $n$  and still admissible, the number of paths expanded can be no larger than uniform-cost.
- Hence the same bounds as uniform-cost apply. (These are worst case bounds). Still exponential unless we have a very good  $h$ !
- In real world problems, we sometimes run out of time and memory.

# Space Problems with A\*

---

- A\* has the same potential space problems as BFS or UCS
- IDA\* - Iterative Deepening A\* is similar to Iterative Deepening Search addresses space issues, but it can require too many depth-first iterations.

# IDA\* - Iterative Deepening A\*

---

Objective: reduce memory requirements for A\*

- Like iterative deepening, but now the “cutoff” is the f-value ( $g+h$ ) rather than the depth
- At each iteration, the cutoff value is the smallest f-value of any path that exceeded the cutoff on the previous iteration
- Avoids overhead associated with keeping a sorted queue of paths, and the open list occupies only linear space.
- Two new parameters:
  - `curBound` (any path with a larger f-value is discarded)
  - `smallestNotExplored` (the smallest f-value for discarded paths in a round)  
when OPEN becomes empty, the search starts a new round with this bound.
  - To make progress on each new depth-first iteration we need to expand all paths with f-value EQUAL to the f-Easier. If any of them reach a goal state then we have found an optimal solution. Otherwise we set the next bound to be the minimum f-value over all paths not added to OPEN because their f-value exceeded `curBound`.

---

# Constructing Heuristics

# Building Heuristics: Relaxed Problem

---

- One useful technique is to consider an easier problem, and let  $h(n)$  be the cost of reaching the goal in the easier problem.
- 8-Puzzle moves.
  - Can move a tile from square A to B if
    - A is adjacent (left, right, above, below) to B
    - **and** B is blank
- Can relax some of these conditions
  1. can move from A to B if A is adjacent to B (ignore whether or not position is blank)
  2. can move from A to B if B is blank (ignore adjacency)
  3. can move from A to B (ignore both conditions).

# Building Heuristics: Relaxed Problem

---

- **#3** “*can move from A to B (ignore both conditions)*”.

leads to the **misplaced tiles** heuristic.

- To solve the puzzle, we need to move each tile into its final position.
- Number of moves = number of misplaced tiles.
- Clearly  $h(n)$  = number of misplaced tiles  $\leq$  the  $h^*(n)$  the cost of an optimal sequence of moves from  $n$ .

- **#1** “*can move from A to B if A is adjacent to B (ignore whether or not position is blank)*”

leads to the **Manhattan distance** heuristic.

- To solve the puzzle we need to slide each tile into its final position.
- We can move vertically or horizontally.
- Number of moves = sum over all of the tiles of the number of vertical and horizontal slides we need to move that tile into place.
- Again  $h(n)$  = sum of the Manhattan distances  $\leq h^*(n)$ 
  - in a real solution we need to move each tile at least that far and we can only move one tile at a time.

# Building Heuristics: Relaxed Problem

---

Comparison of IDS and A\* (average total paths expanded ):

Depth	IDS	A*(Misplaced) <b>h1</b>	A*(Manhattan) <b>h2</b>
10	47,127	93	39
14	3,473,941	539	113
24	---	39,135	1,641

Let **h1**=Misplaced, **h2**=Manhattan

- Does **h2 always** expand fewer paths than **h1**?
  - Yes! Note that **h2 dominates h1**, i.e. for all  $n$ :  $h1(n) \leq h2(n)$ . From this you can prove **h2** is better than **h1** (once both are admissible).
  - Therefore, among several admissible heuristic the one with highest value is better (will cause fewer paths to be expanded).

# Building Heuristics: Relaxed Problem

---

The **optimal** cost of reaching the goal in the relaxed problem is an **admissible heuristic** for the original problem!

**Proof Idea:** the optimal solution in the original problem is a solution for relaxed problem, therefore it must be at least as expensive as the optimal solution in the relaxed problem.

So admissible heuristics can sometimes be constructed by finding a relaxation whose optimal solution can be easily computed.



# Building Heuristics: Relaxed Problem

---

The Euclidean distance in the Romanian Travel state space is another example of a using a relaxed problem to compute a heuristic.

In this case the relaxed problem is one where there is a direct straight line road from every city to the goal city. In this relaxed problem the goal can always be **optimally** reached from any state by taking the direct road. Hence the Euclidean distance is the optimal solution cost in the relaxed problem and thus it is an **admissible heuristic** in the real problem.

# Building Heuristics: Pattern databases

---

- Try to generate admissible heuristics by solving a subproblem and storing the exact solution cost for that subproblem
- See Chapter 3.6.3 if you are interested.

*	2	4
*		*
*	3	1

Start State

	1	2
3	4	*
*	*	*

Goal State

# Building Heuristics: Counting Necessary Actions.

---

- In some problems we can find for every state a set of necessary actions...actions that must be executed to reach the goal. Summing their costs yields an admissible heuristic.
- E.g., moving goods with trucks.
  - We know that every good not at its goal location must be loaded onto some truck and later unloaded.
  - We know that every good must also be moved in a truck at least the Euclidean distance from its current location to its goal location. **However, we have to be careful.** For an admissible heuristic we cannot double count these distances. That is, one move of the truck will move all of the goods in it, we cannot sum the moving distance of different goods in the same truck.